

CPSC 213

Introduction to Computer Systems

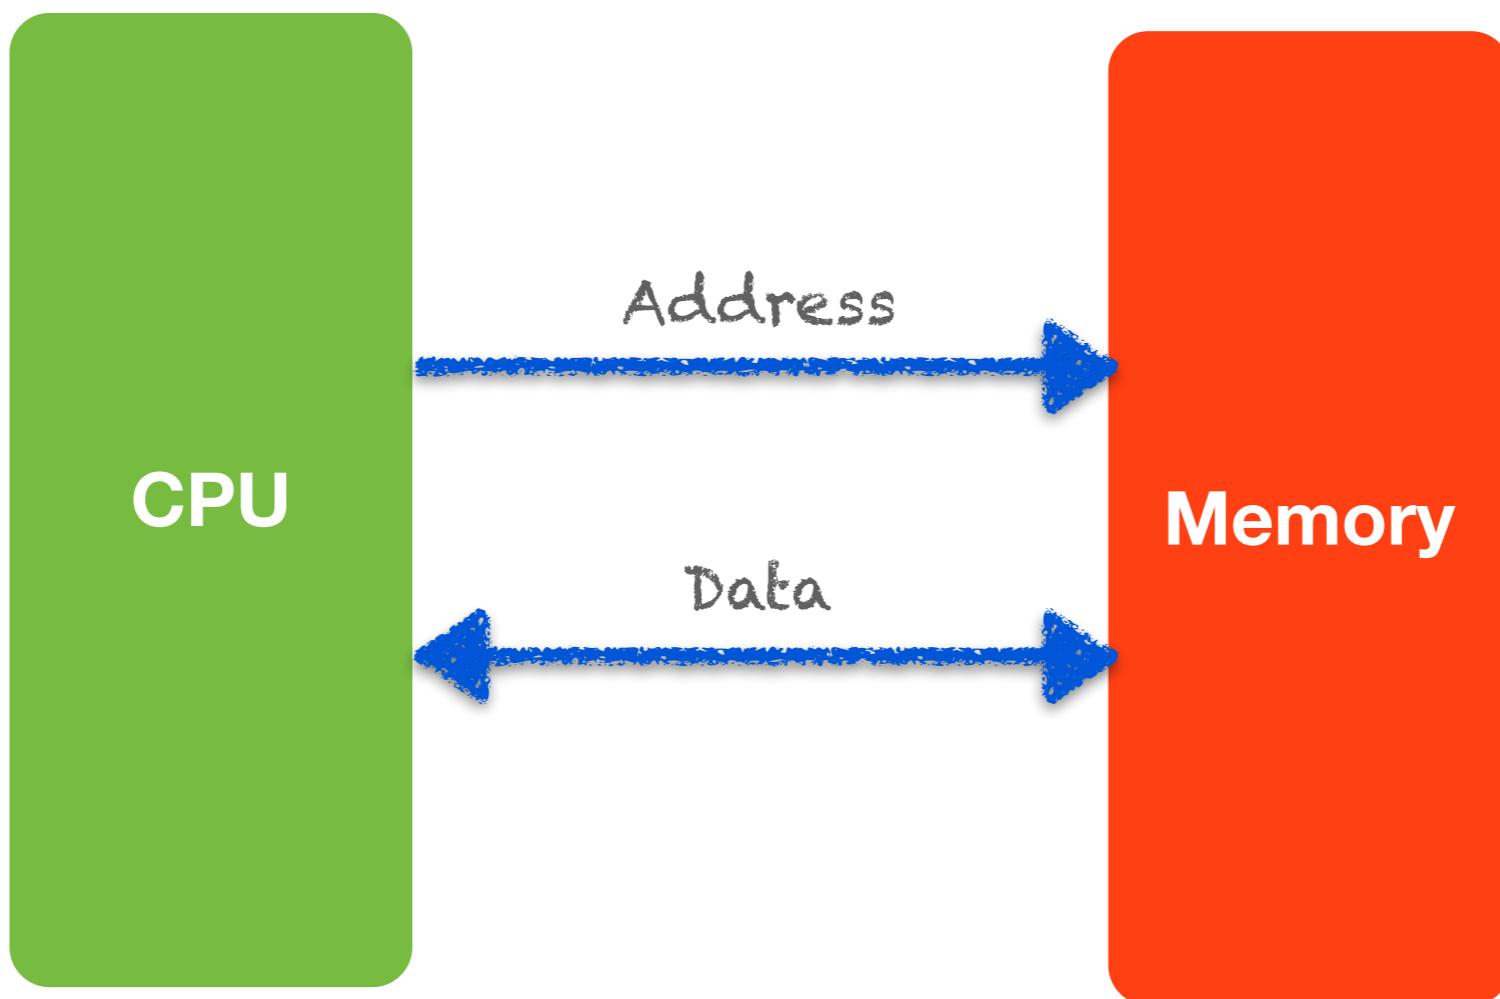
Summer Session 2019, Term 2

Unit 1b – July 4, 9

Static Scalars and Arrays

Recap

- ▶ Memory = big bag of bytes
- ▶ Program interprets those bytes!
- ▶ Multi-byte integers and endianness
- ▶ Bit operations



Overview of Unit 1b

▶ Reading

- Companion: 1, 2-1-2.3, 2.4.1-2.4.3
- Textbook: 3.1-3.2.1

▶ Reference (as needed)

- Textbook: 3.1-3.5, 3.8, 3.9.3

▶ Learning Objectives

- list the basic components of a simple computer and describe their function
- describe ALU functionality in terms of its inputs and outputs
- describe the exchange of data between ALU, registers, and memory
- identify and describe the basic components of a machine instruction
- outline the steps a RISC machine performs to do arithmetic on numbers stored in memory
- translate between array-element offsets and indices.
- distinguish static and dynamic computation for access to global scalars and arrays in C
- describe the tradeoffs between static and dynamic arrays
- describe why C does not perform array-bounds checking and what that means for C programs
- translate between C and assembly language for access to global scalars and arrays
- translate between C and assembly for code that performs simple arithmetic
- explain the difference between arrays in C and Java
- use C's dereference and address-of operators and pointer arithmetic to access elements of arrays

Our Approach

▶ Develop a model of computation

- that is rooted in what the machine actually does
- by examining C, bit-by-bit (comparing to Java as we go)

▶ The processor

- we will design (and you will implement) a simple instruction set
- based on what we need to compute C programs
- similar to a real instruction set (MIPS)

▶ The language

- we will act as compiler to translate C into machine language
- bit by bit, then putting the bits together to do interesting things
- edit, debug and run using simulated processor to *visualize* execution

We will start here ...

Java:

```
public class Foo {  
    static int a;  
    static int[] b; // array is not static, so skip for now  
  
    public void foo () {  
        a = 0;  
    }  
}
```

C:

```
int a;  
int b[10];  
  
void foo () {  
    a = 0;  
    b[a] = a;  
}
```

The CPU

- ▶ CPUs execute *instructions*, not C or Java code
- ▶ Execution proceeds in three phases:
 - Fetch - load the next instruction from memory
 - Decode - figure out what the instruction will do
 - Execute - do what the instruction asks
- ▶ These phases are looped over and over again forever

The CPU

- ▶ Instructions are very simple

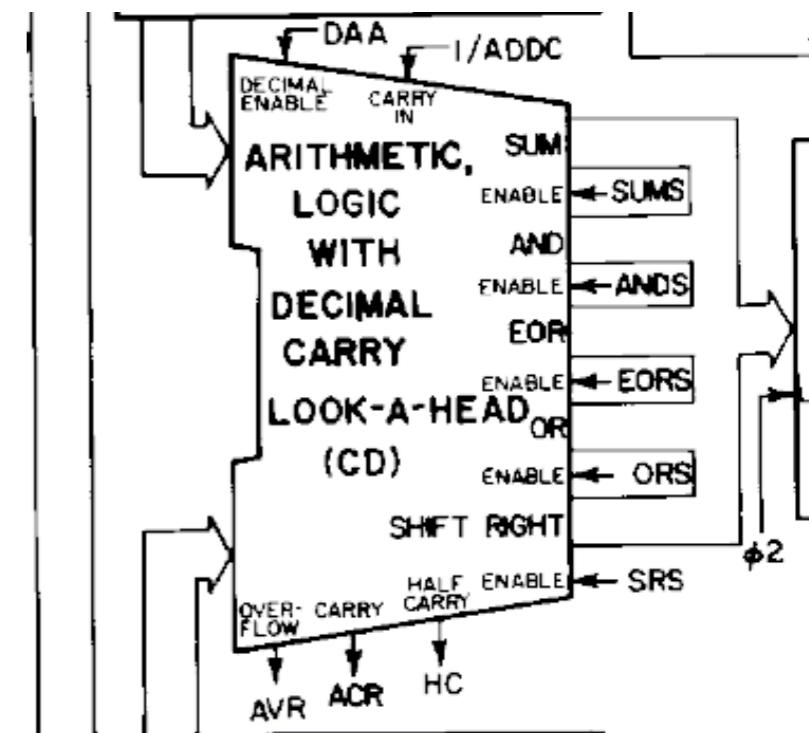
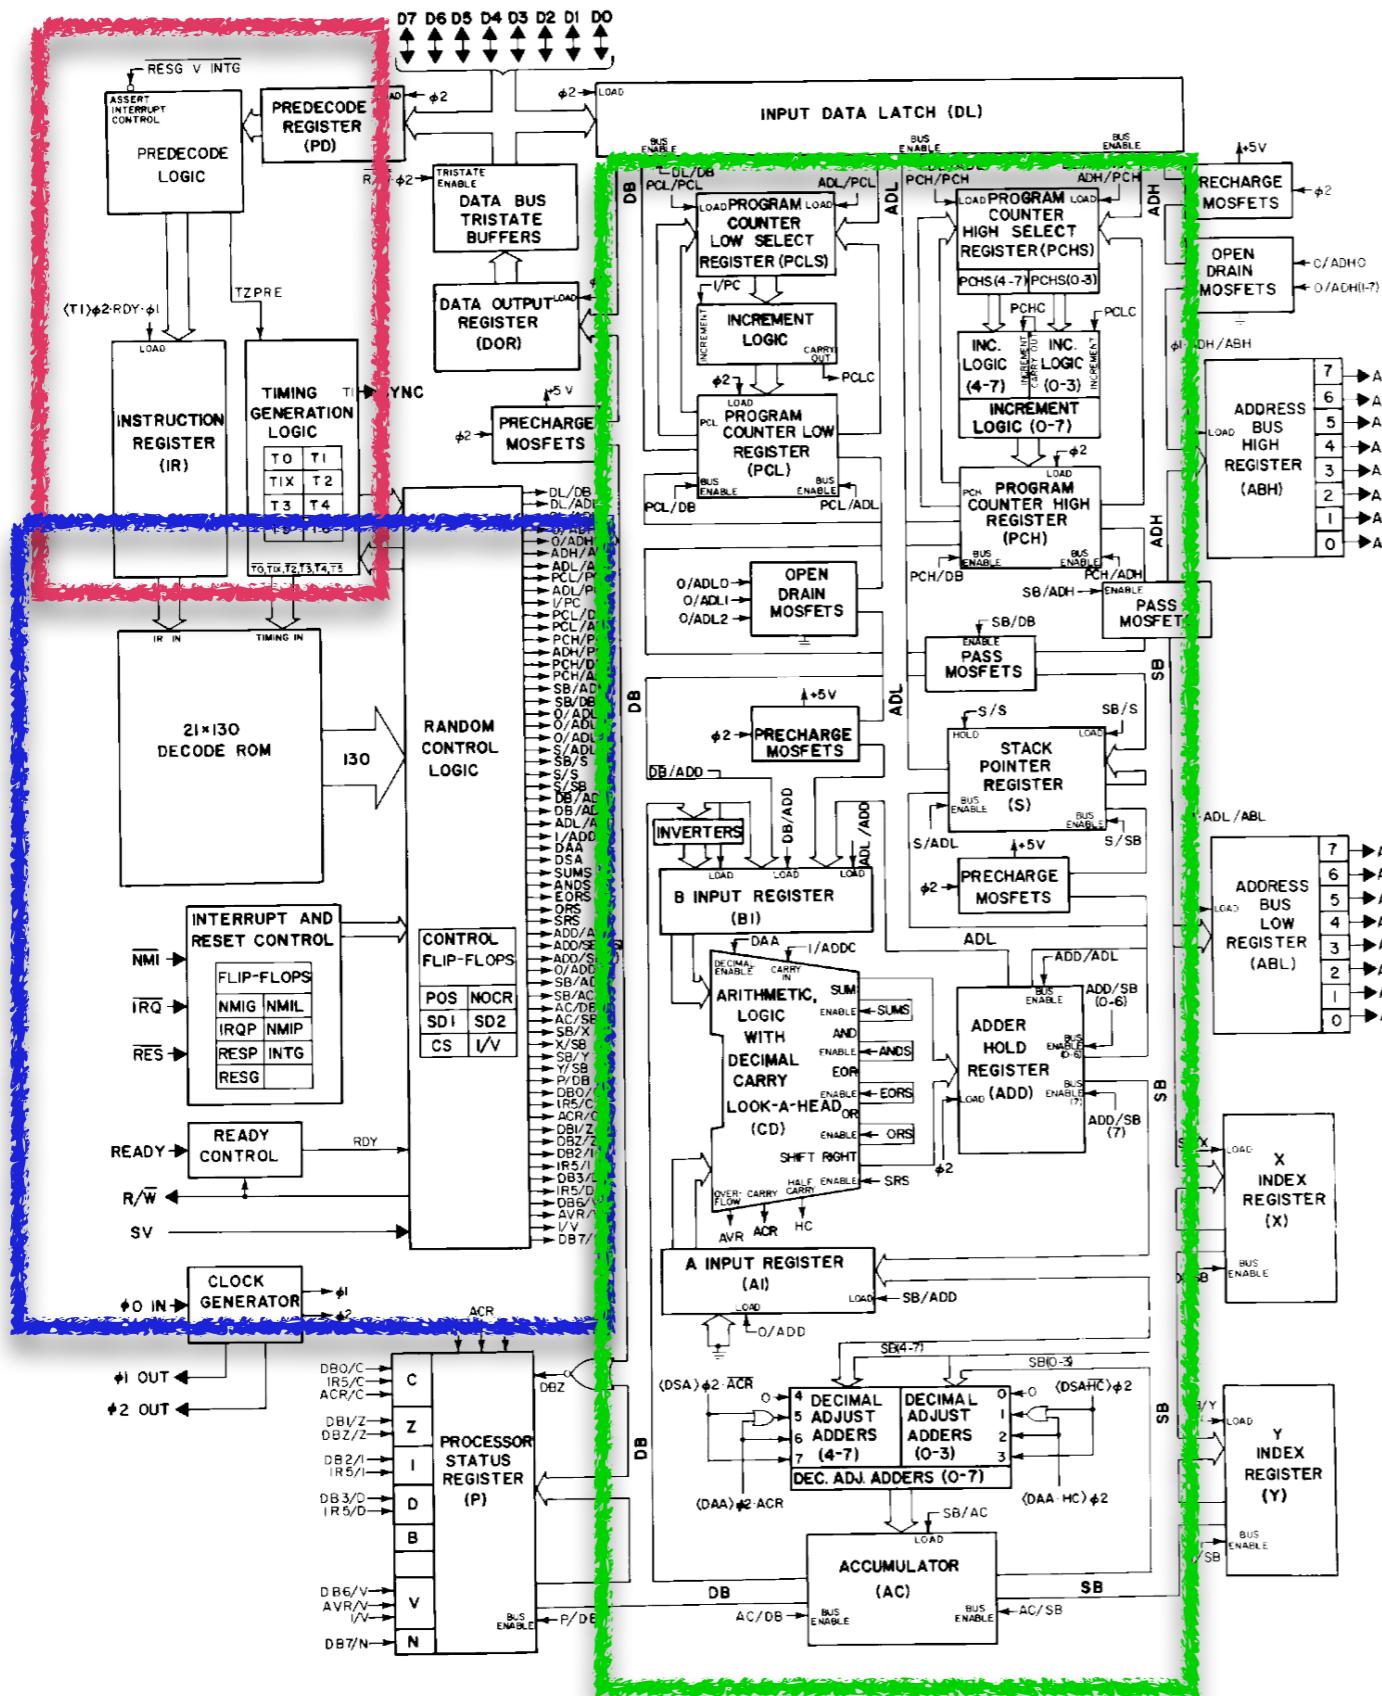
- Add/Subtract two numbers
- And/Or two numbers
- Bit shift a number
- Read/Write memory
- Control flow (later...)

- ▶ Operations carried out by the ALU (Arithmetic & Logic Unit)

- ▶ ALU is programmed by the Execute step

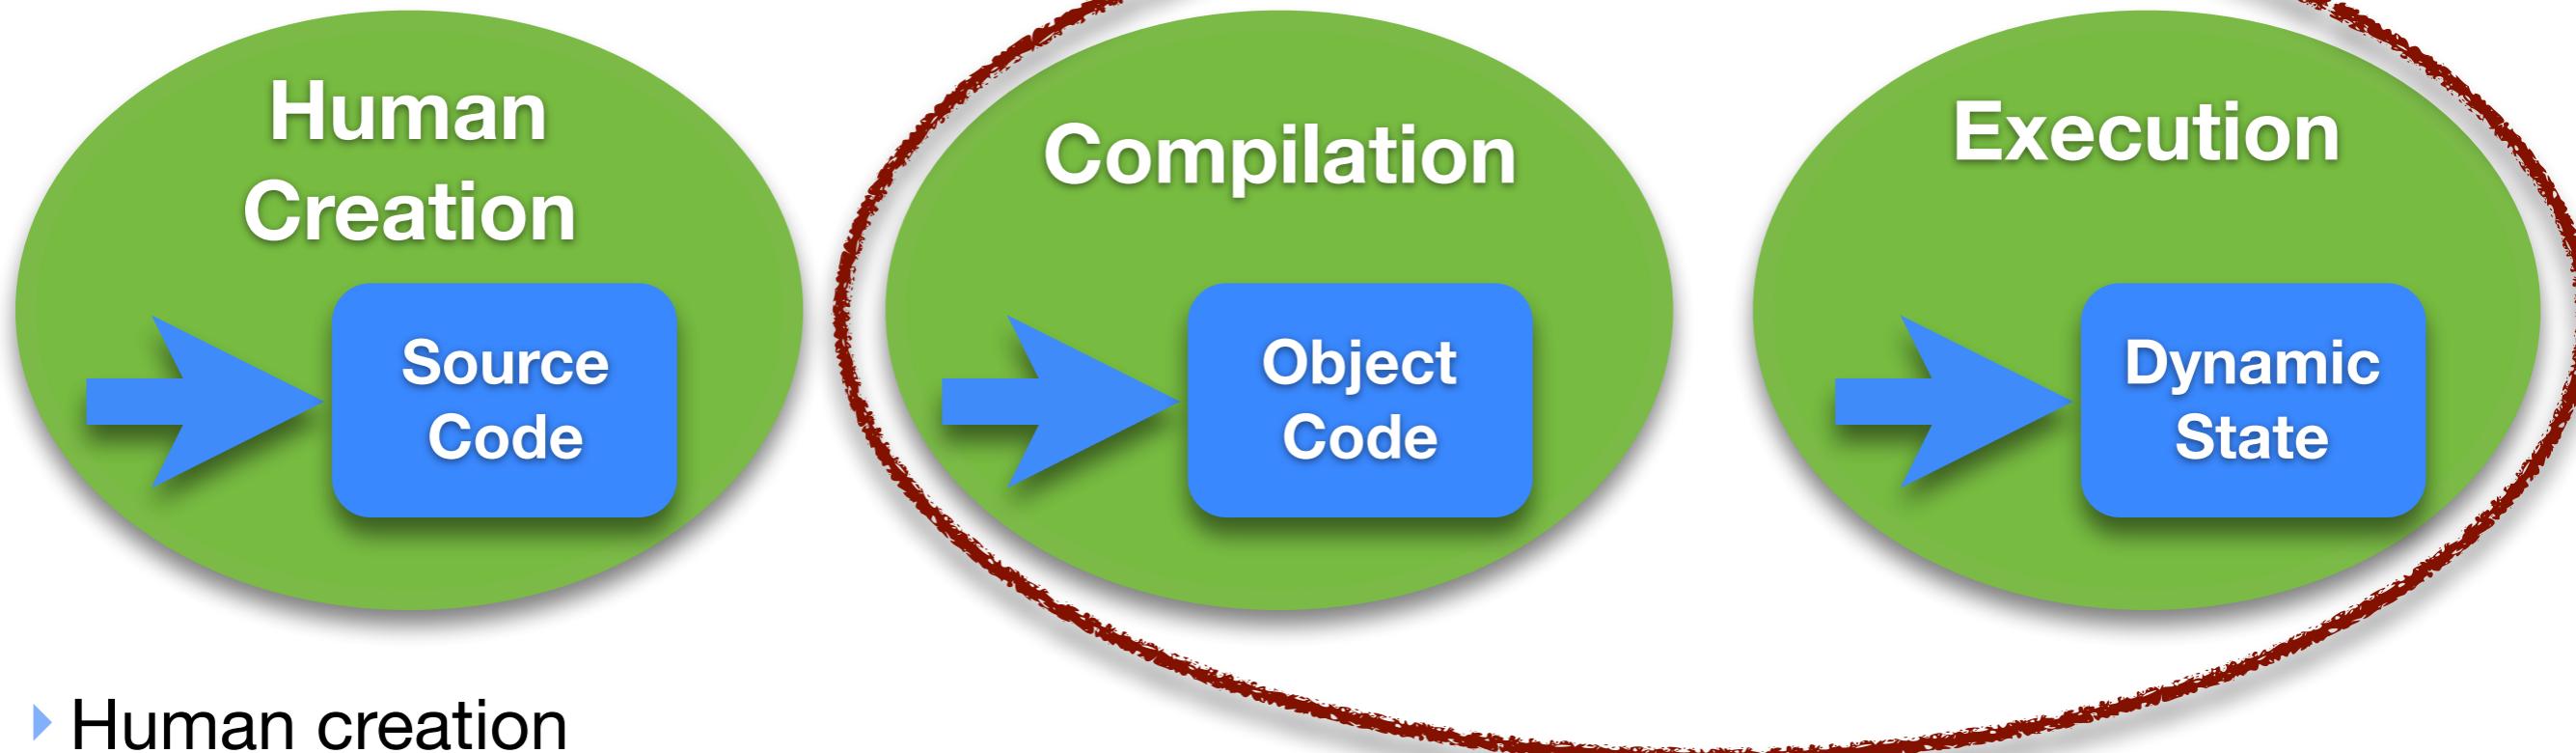
- ALU carries out one mathematical or logical operation based on inputs

The CPU



Donald F. Hanson, WCAE 1995

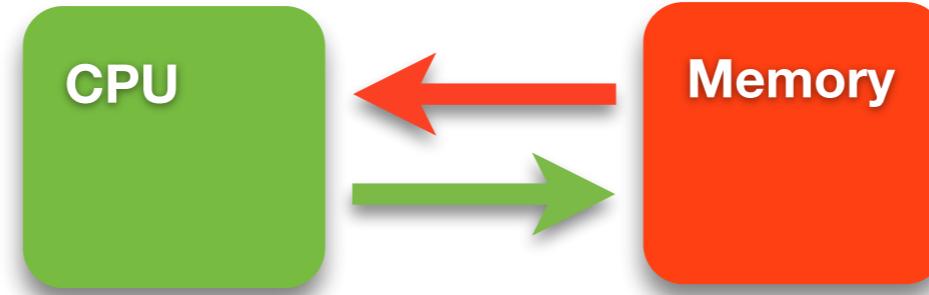
Phases of Computation



- ▶ Human creation
 - design program and describe it in high-level language
- ▶ Compilation
 - convert high-level, human description into machine-executable text
- ▶ Execution
 - a physical machine executes the text
 - parameterized by input values that are unknown at compilation
 - producing output values that are unknowable at compilation
- ▶ Two Crucial Phases of Computation
 - **STATIC** anything that the compiler can possibly compute is called *static*
 - **DYNAMIC** anything that can not possibly be known until the program runs is called *dynamic*

First An Introduction to Computing with a Simple Machine

The Processor (CPU)



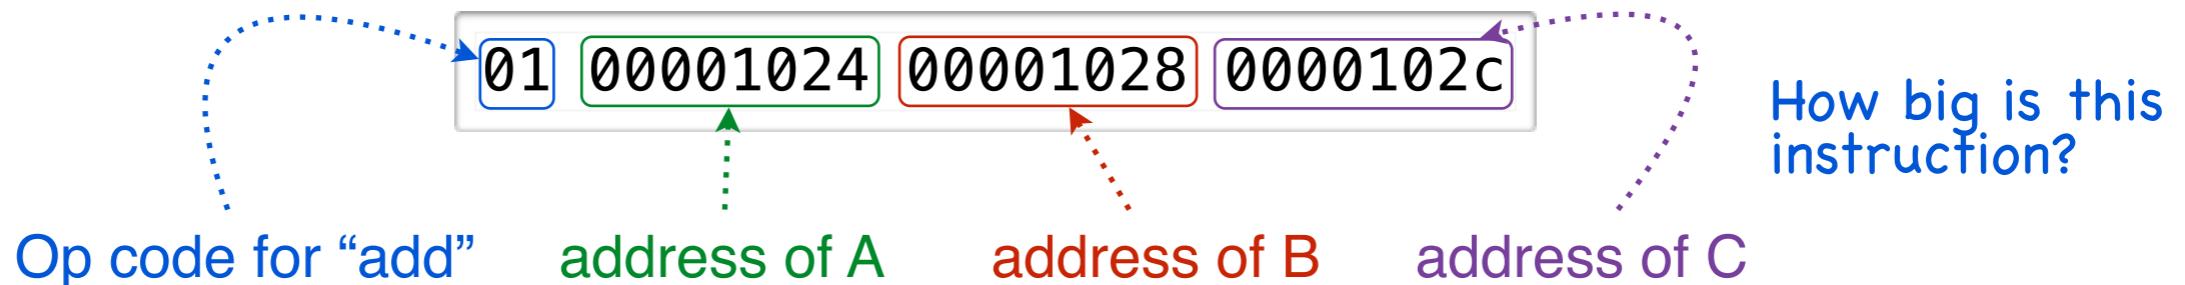
How many memory accesses?

► Implements a set of simple instructions

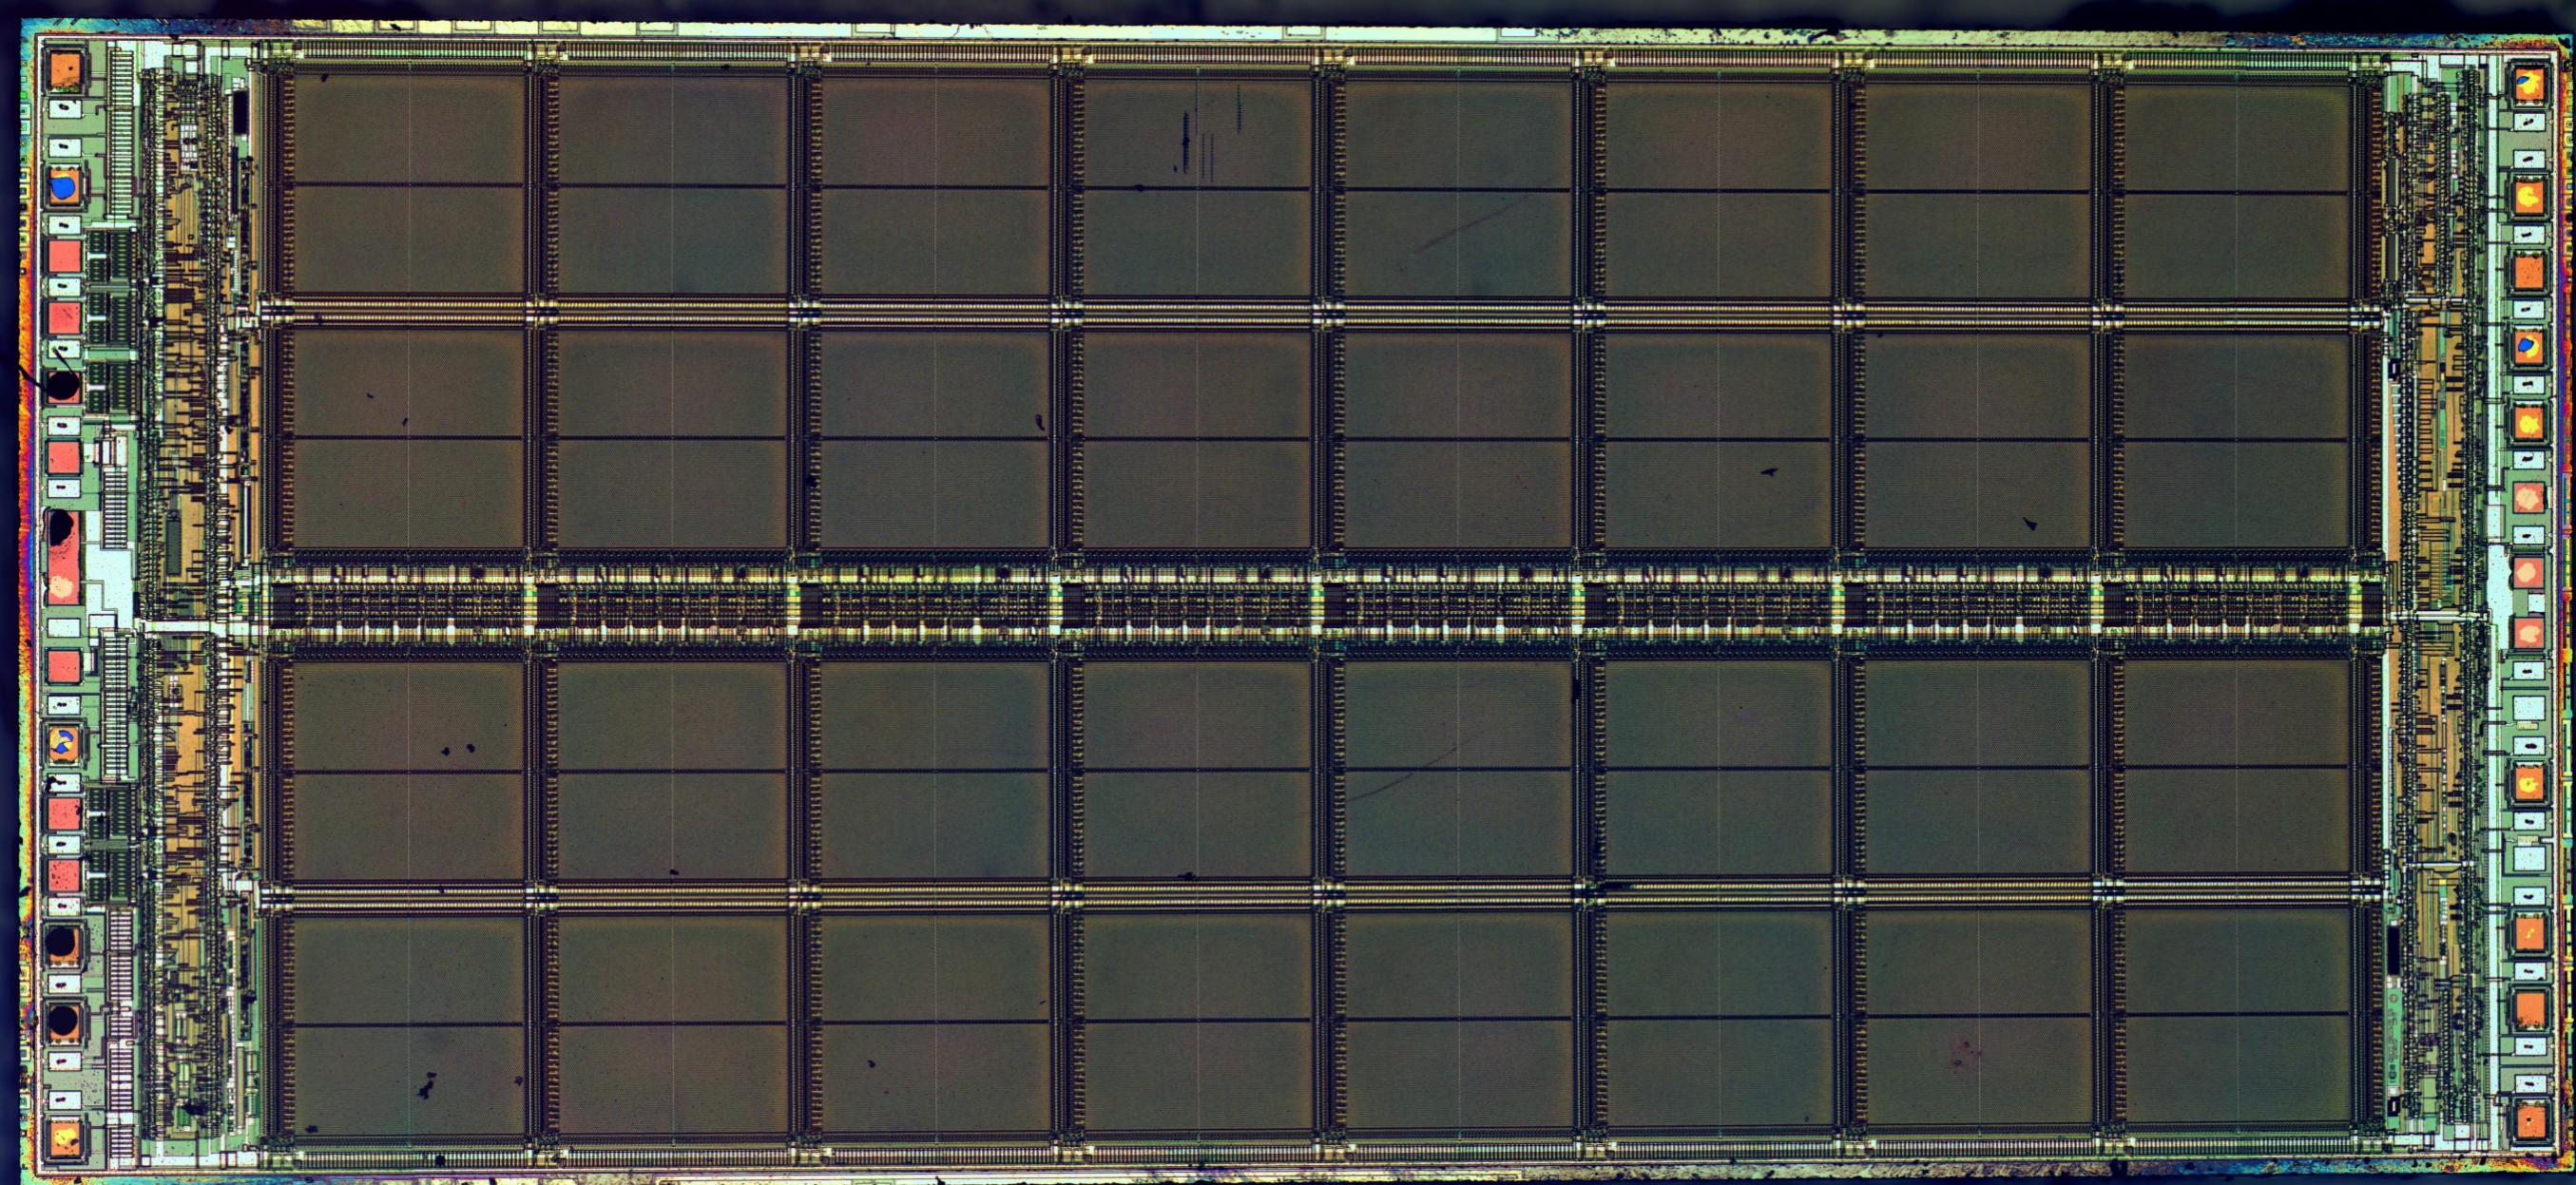
- each instruction is implemented using logic gates, built from transistors
- these transistors are the fundamental mechanism of computation
- the fewer and simpler the better

► For example, we'll want to do arithmetic

- we'll have an add instruction that does something like: $C = A + B$
 - A, B and C are values stored in memory – named by 32-bit addresses (in our case, 64-bits typically today)
- every instruction is encoded as a set of bits stored in memory
 - the instruction will need to name the operation and the three operands (assume their addresses are constants)
- the instruction might look something like this (in hex)



The Memory (RAM)



Micron MT4C1024 128KB RAM, zeptobars.org

The Problems with Memory Access



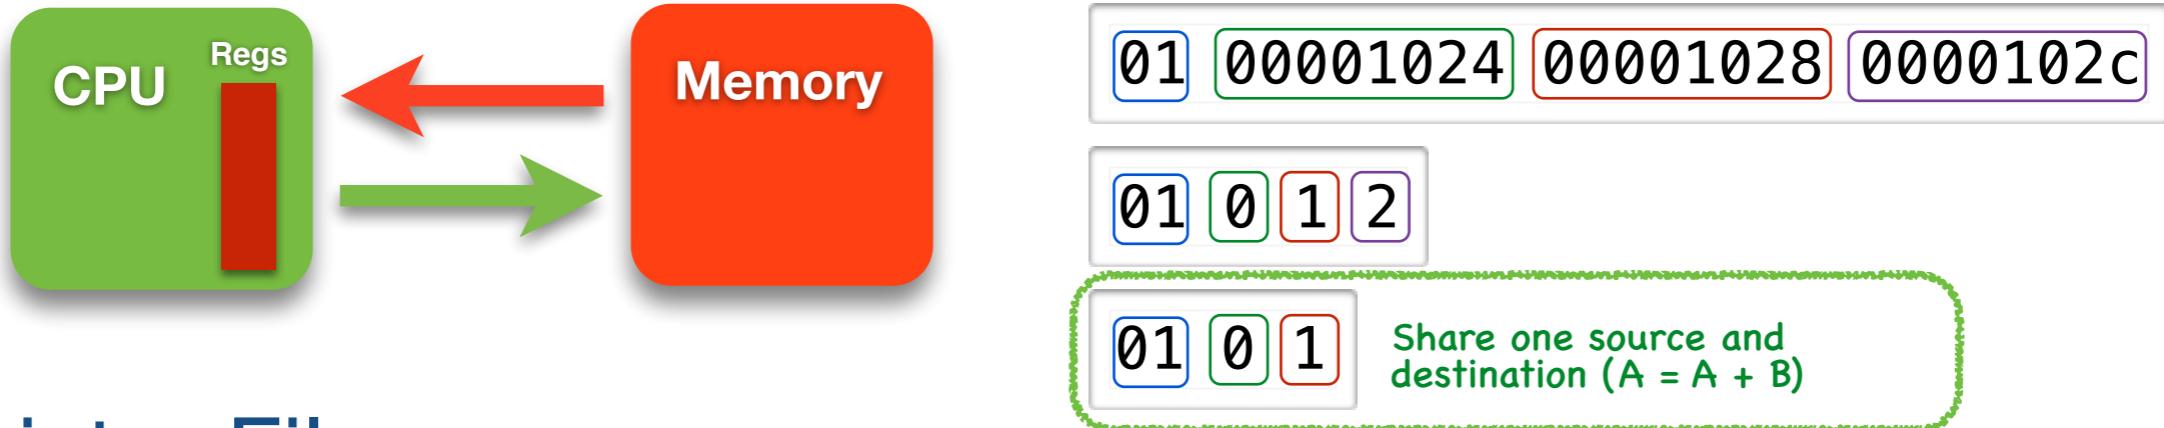
▶ Accessing memory is slow

- ~20–100 cycles for every memory access
 - there are tricks to help (more in CPSC 313), but it's still slow
- fast programs avoid accessing memory whenever possible

▶ Big instructions are costly

- memory addresses are big, thus so are instructions that access memory
- if most instructions are big, then programs will be big too
 - memory is a finite resource; we should use it efficiently
- if most instructions are big, then programs will be slow
 - CPU reads an instruction every cycle, caching them locally in case they are accessed again
 - if enough instructions are cached, the CPU rarely waits to get instruction from memory
 - if instructions are big, fewer will fit in the cache and there will be more memory stalls
 - the CPU must stall whenever it needs an instruction that is not stored locally ... to get it from memory
 - also, big instructions take longer to transfer to CPU and use more resources in the process

The General-Purpose Register File



▶ Register File

- small, fast, on-chip memory directly manipulated by instructions
- each register is named by a number (e.g., 0 – 7) 8 times smaller than a 32-bit memory address

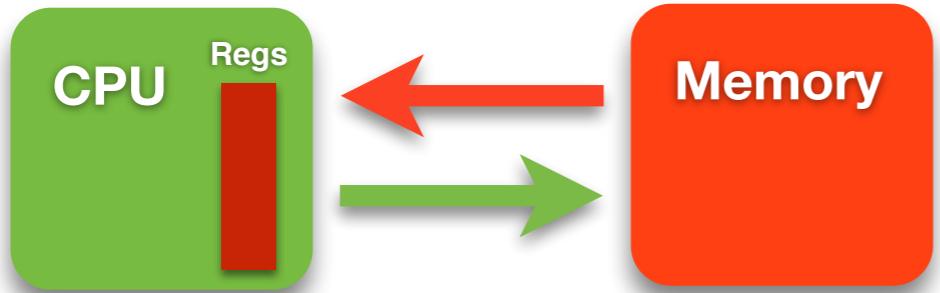
▶ Registers

- size of largest integer (e.g., 32 bits or maybe 64)
- built from fast memory (roughly single-cycle access)

▶ Instructions

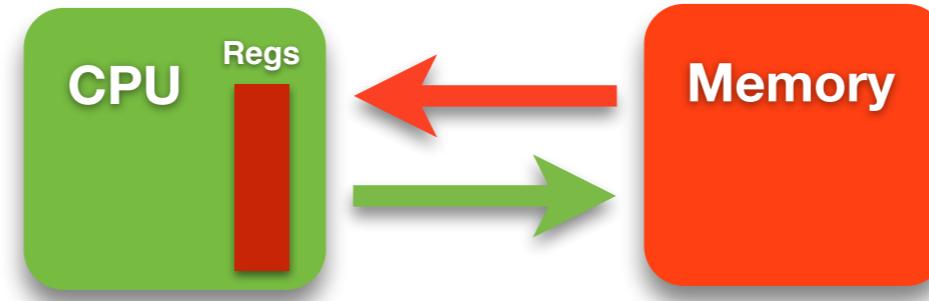
- load data from memory into register; store data from register into memory
 - these are large and slow
- other instructions access data in registers
 - these are small and fast

Special-Purpose Registers



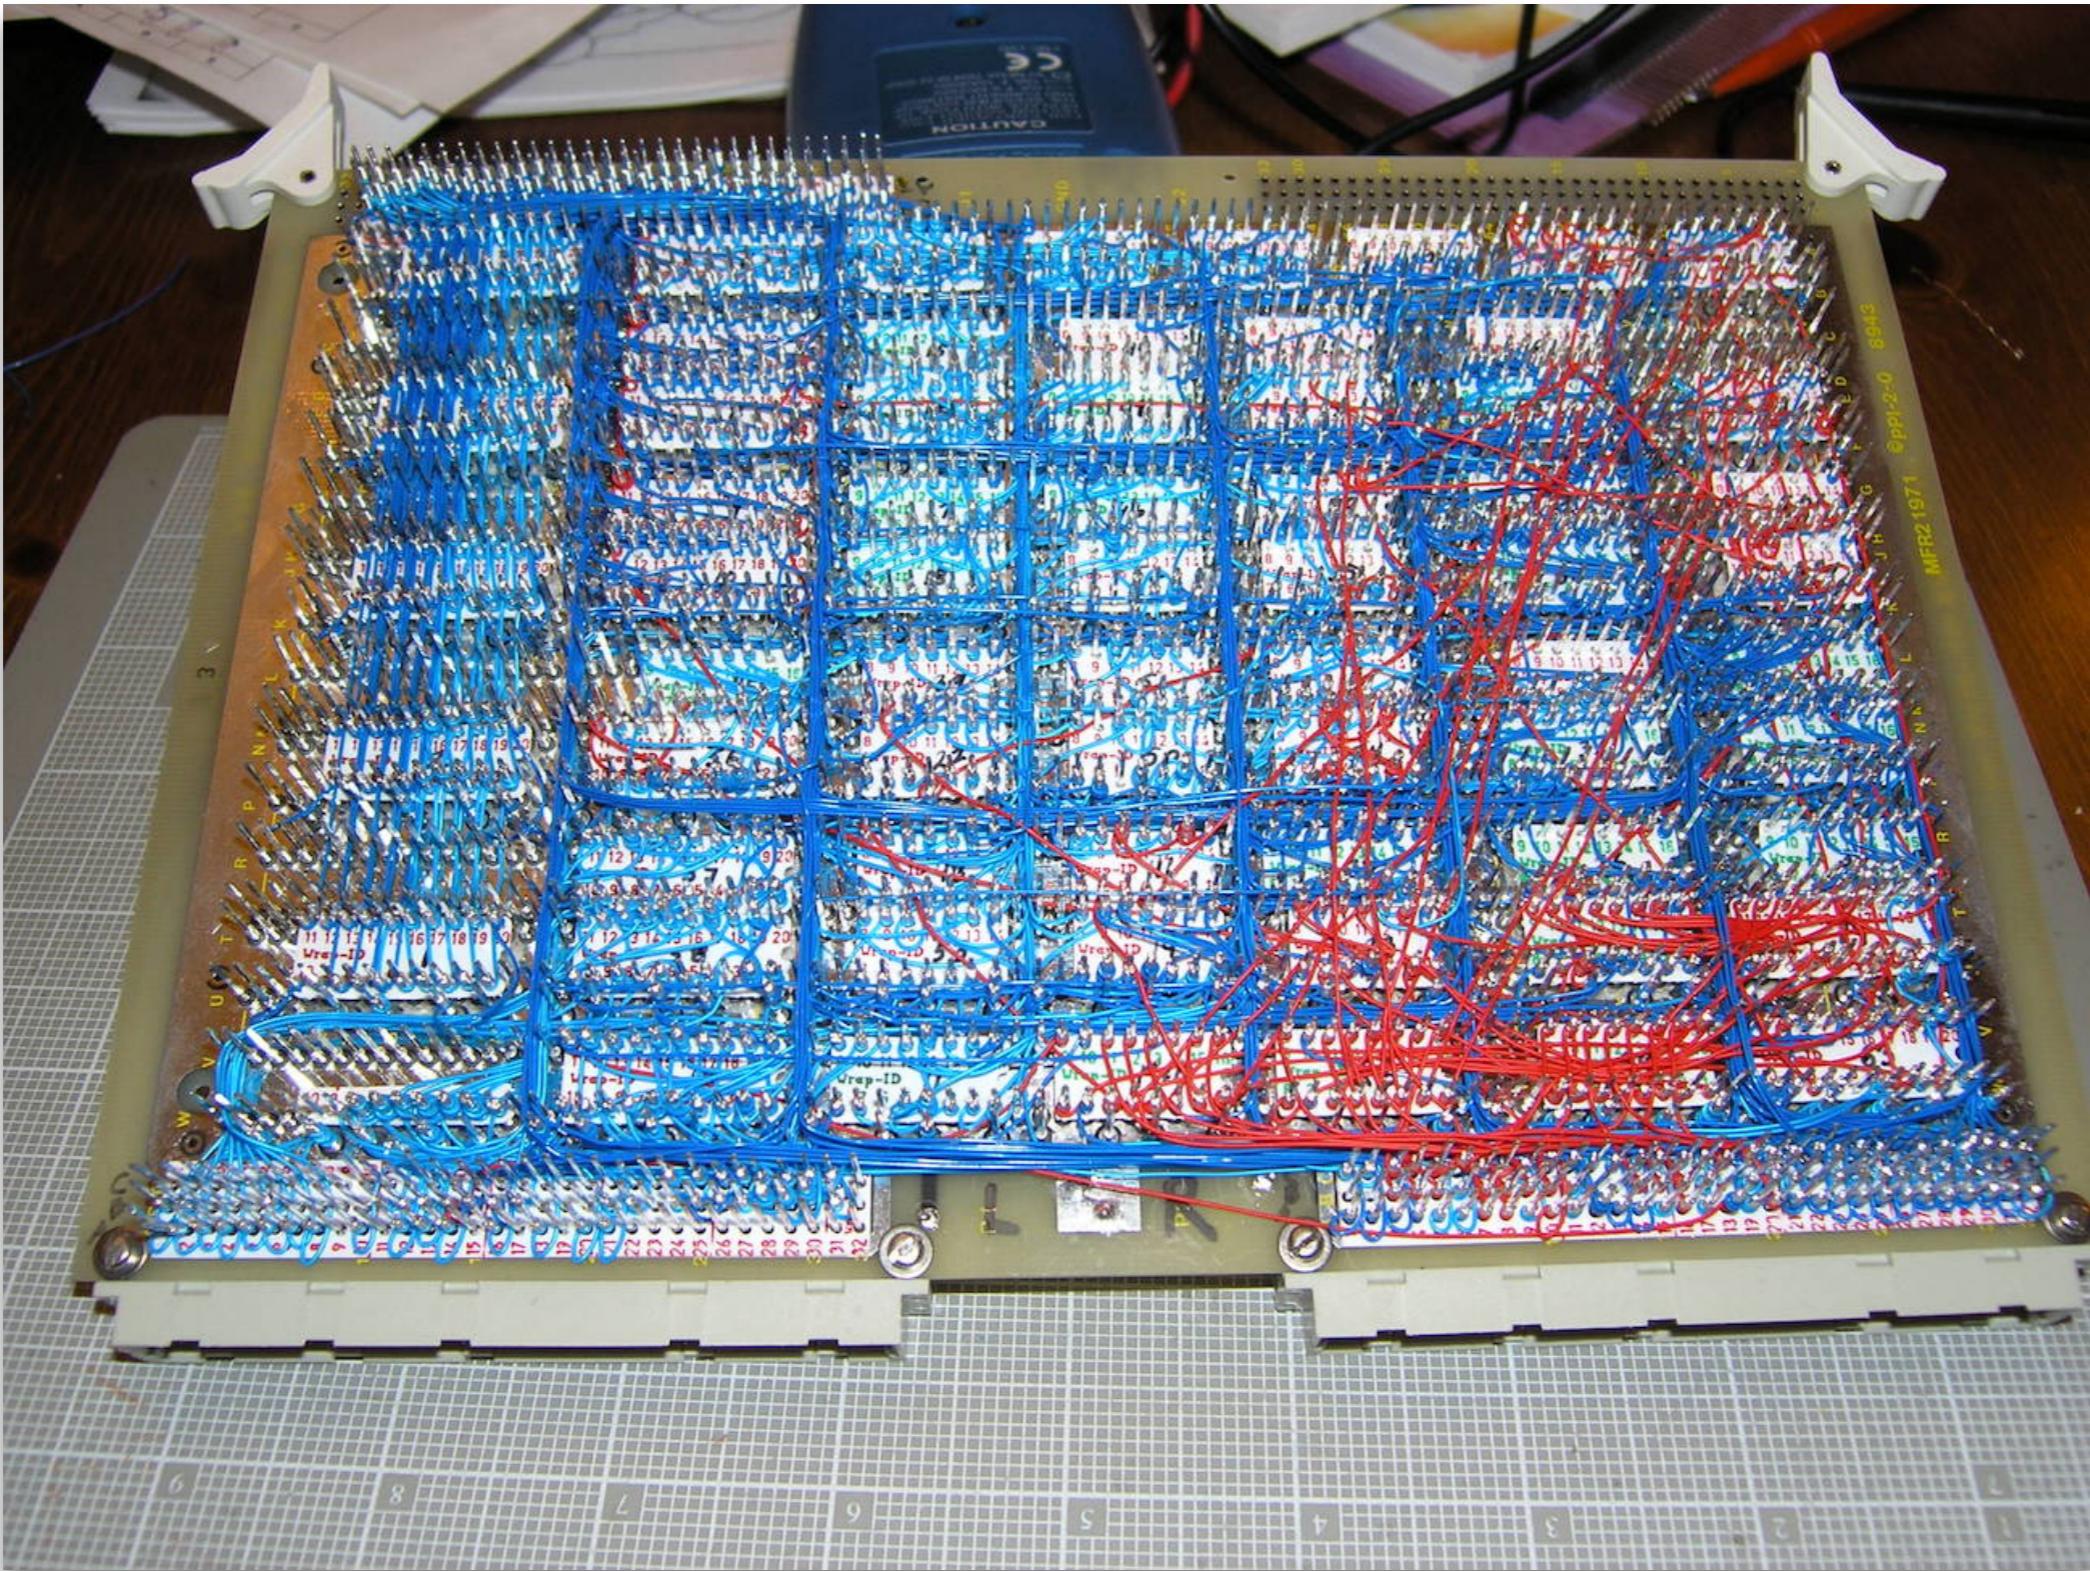
- ▶ A *special-purpose register* can only be used for certain purposes
 - May not be accessible by all instructions
 - May have special meaning or be treated specially by CPU
- ▶ Some CPUs have many
- ▶ **PC, the Program Counter**
 - Address of the next instruction to execute
- ▶ **IR, the Instruction Register**
 - Instruction that has just been loaded from RAM
 - Separate smaller registers for the parts

Simple Machine Hardware Architecture



- ▶ **CPU** The Program and the Machine
- ▶ Register File
- ▶ Main Memory

Let's build a CPU!



Home-made CPU control board, homebrewcpu.com

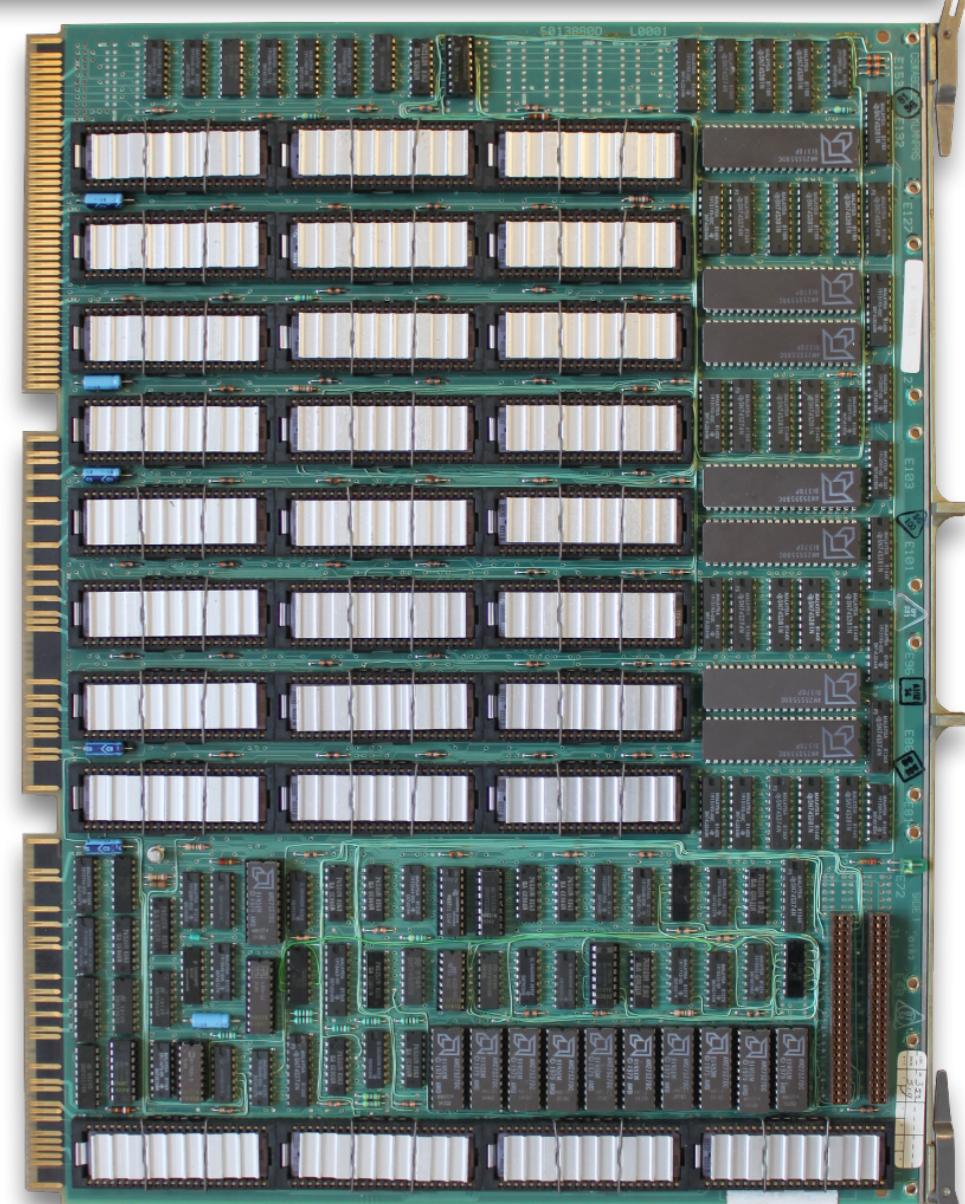
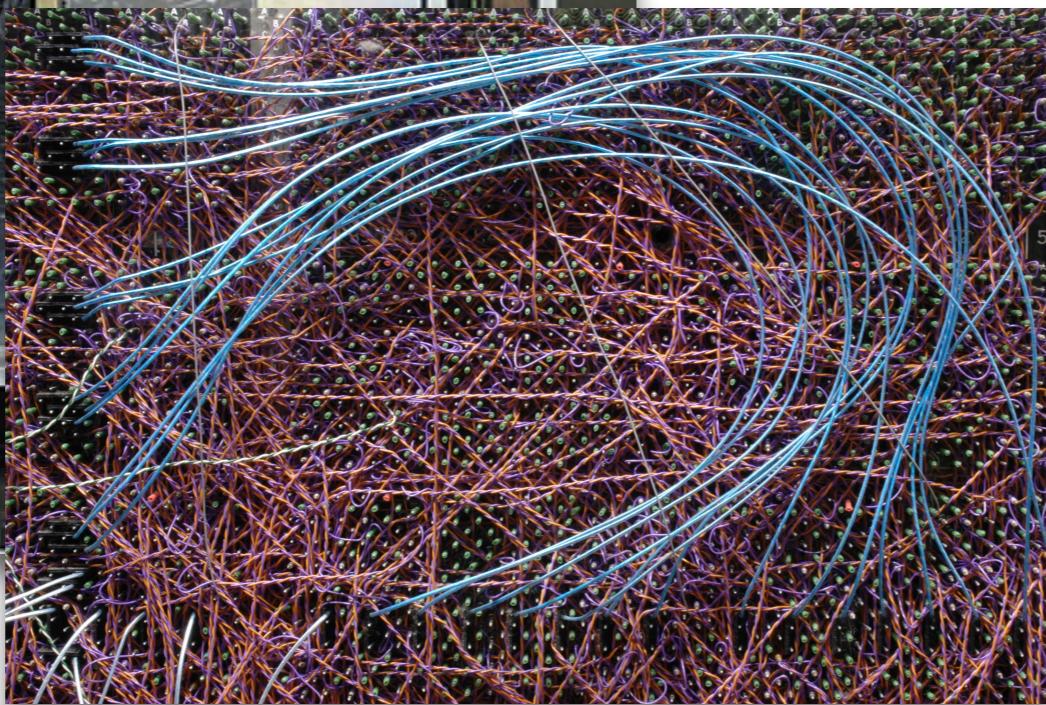
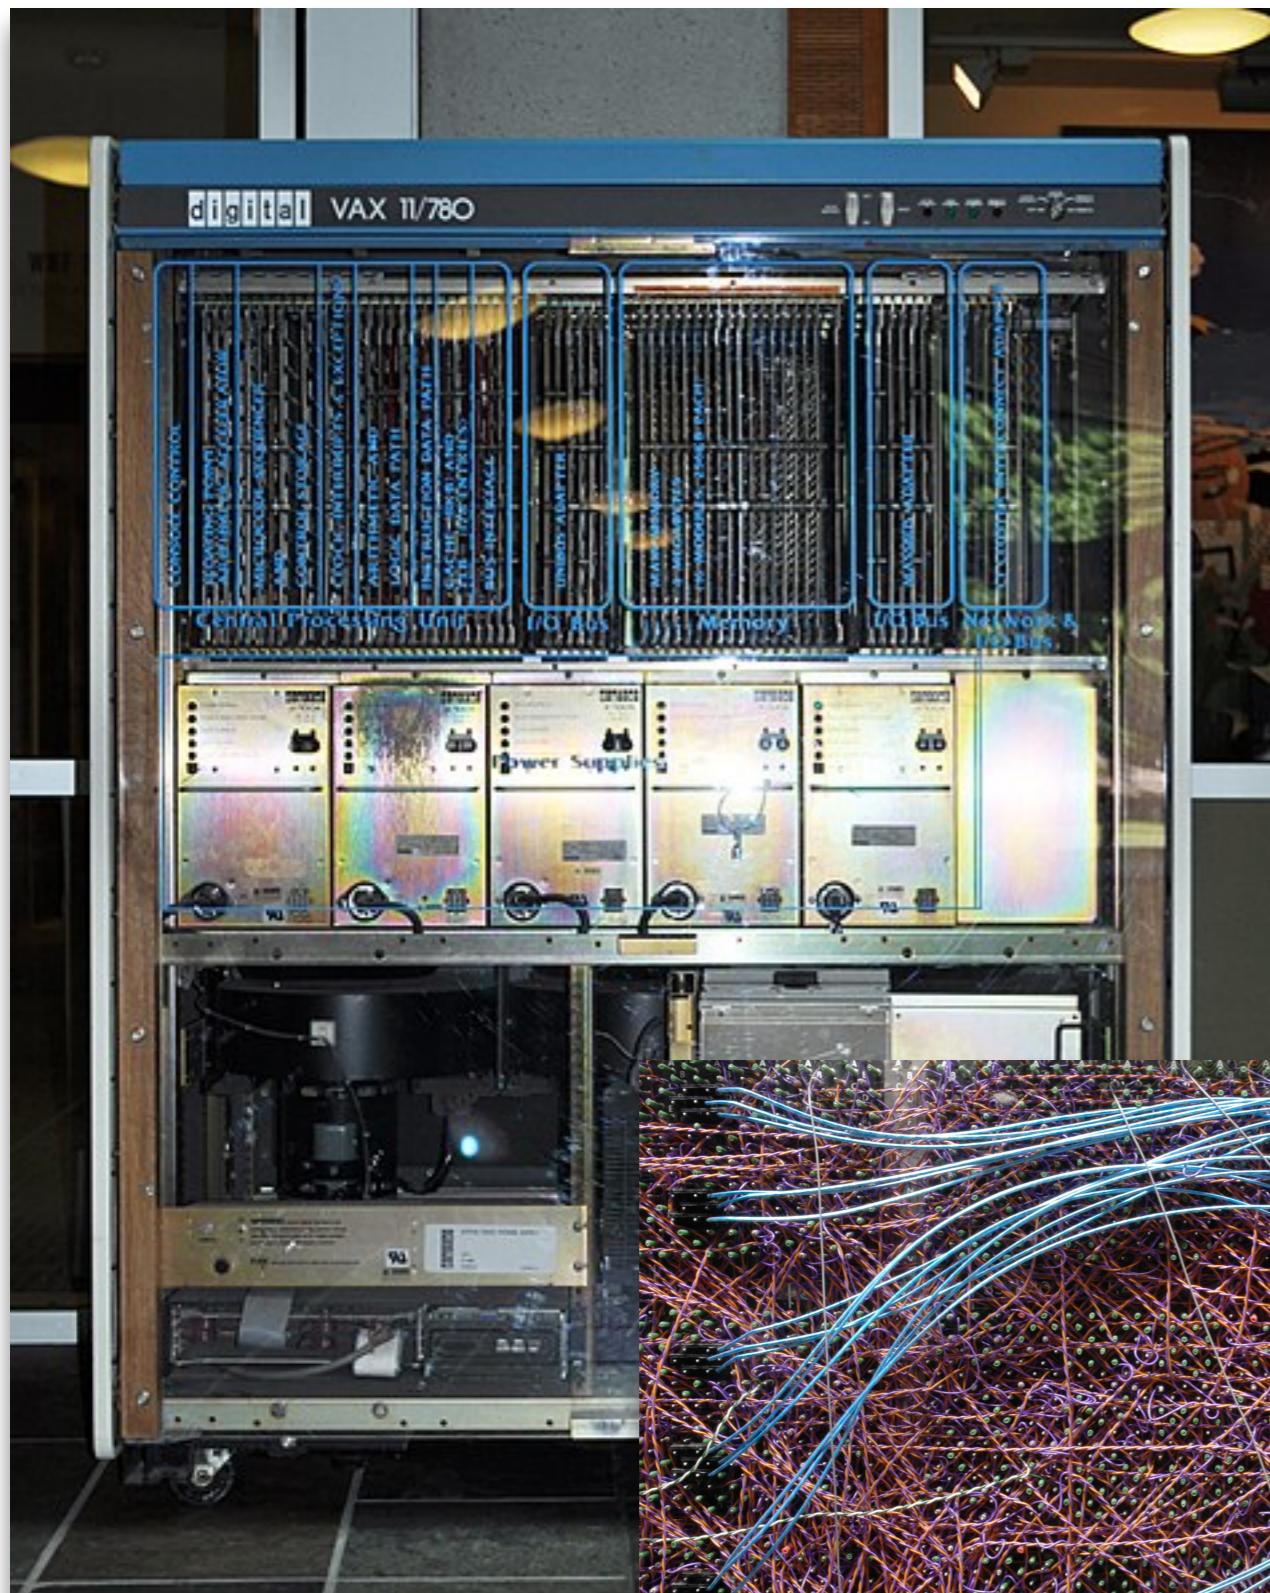
The Simulator ...

The Processor

Instruction Set Architecture (ISA)

- ▶ The ISA is the interface to a processor implementation
 - defines the instructions the processor implements
 - defines the format of each instruction
- ▶ Types of instruction
 - math
 - memory access
 - control transfer (gotos and conditional gotos)
- ▶ Design alternatives
 - simplify compiler design (CISC such as Intel Architecture 32 (aka x86))
 - simplify processor implementation (RISC)
- ▶ Instruction format
 - is a sequence of bits (or, typically, in hex)
 - an opcode and set of operand values
- ▶ Assembly language
 - symbolic representation of machine code

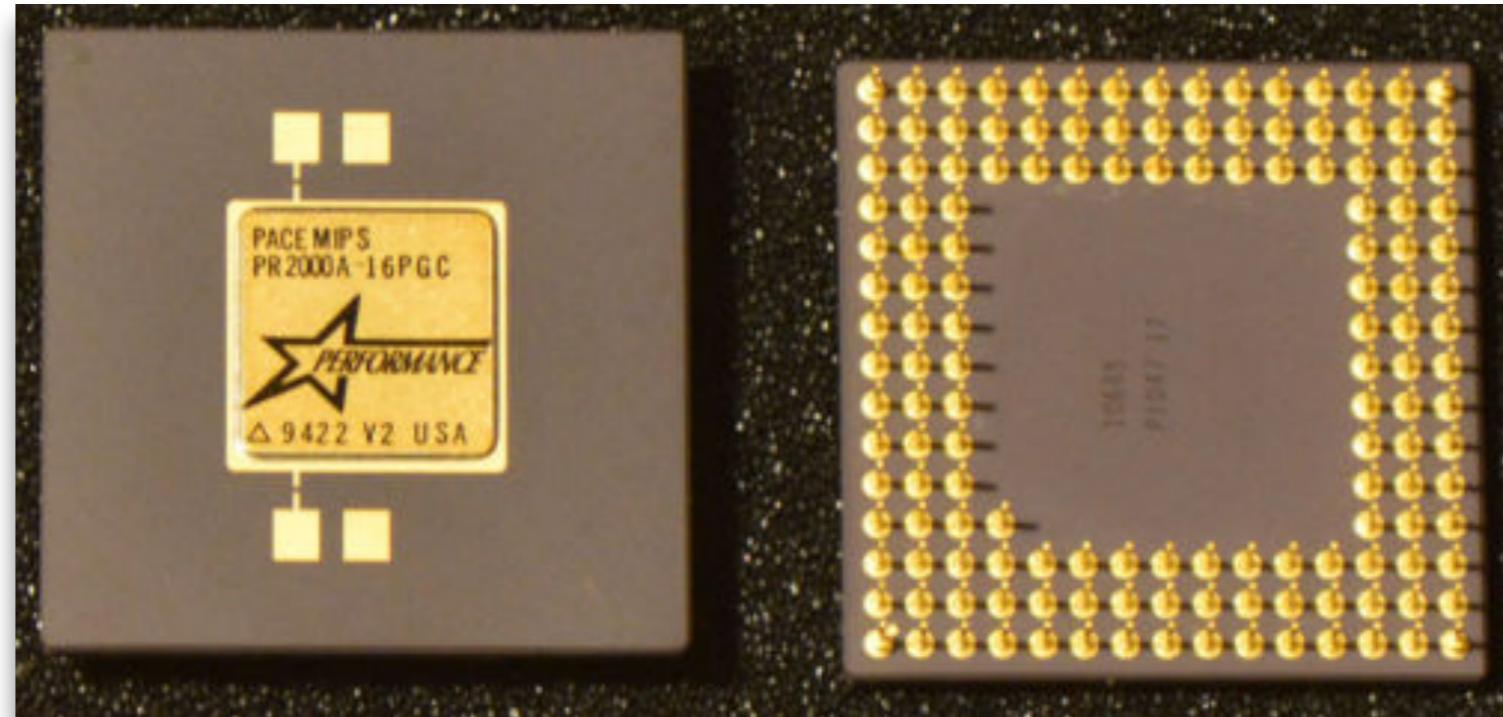
DEC VAX 11/780



MIPS R2000

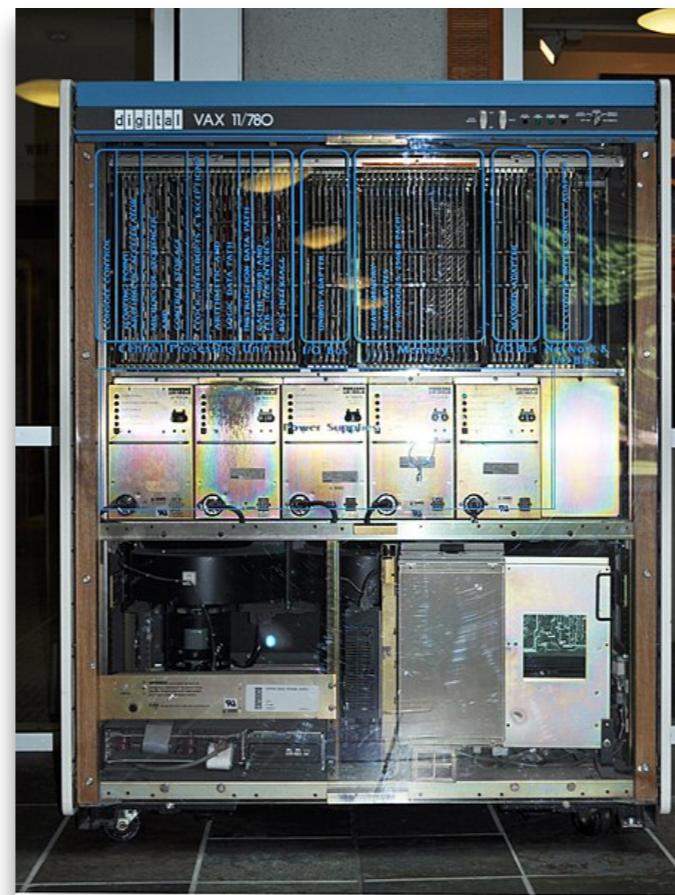
▶ R2000

- 1986
- 8.3 to 15 Mhz
- 110,000 transistors
- 80 mm² die
- 2,000 nm feature size CMOS

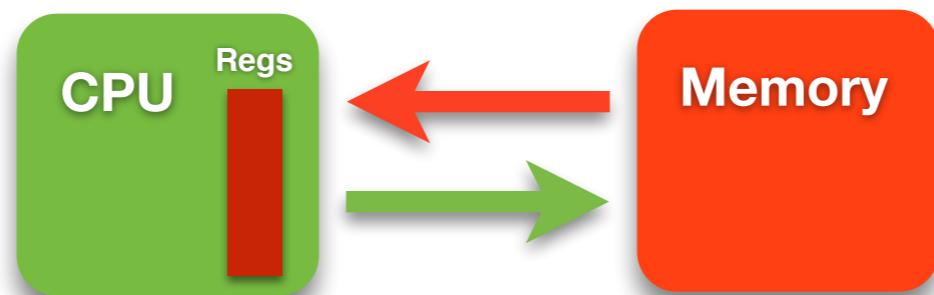


▶ VAX 11/780

- 1977
- 5 Mhz
- 2-MB – 8-MB main memory



Describing Instruction Semantics



▶ RTL is

- a simple, convenient pseudo language to describe instruction semantics
- easy to read and write, directly translated to machine steps

▶ Syntax

- each line is of the form **LHS** \leftarrow **RHS**
- **LHS** is memory or register specification
- **RHS** is constant, memory, or arithmetic expression on two registers

▶ Register and Memory are treated as arrays

- $m[a]$ is value in memory at address **a**
- $r[i]$ is register number **i**

▶ For example

- $r[0] \leftarrow 10$
- $r[1] \leftarrow m[r[0]]$
- $r[2] \leftarrow r[0] + r[1]$

Static Variables of Built-In Types

Variables

▶ What are they?

- named storage location for values

▶ What features do they have?

- name
- size
- type
- scope
- lifetime
- memory location; i.e., an address
- value

▶ Static or Dynamic in Java (or if you know it, C/C++)?

- which of these are determined at compile time?
- which are determined (and thus can change) while the program is running?

▶ Allocation

- is assigning a memory location (i.e., an address) to a variable
- WHEN does this happen? HOW is location found?

Static Variables, Built-In Types (S1-global-static)

▶ Java

- static data members are allocated to a class, not an object
- they can store built-in scalar types or references to arrays or objects (references later)

```
public class Foo {  
    static int a;  
    static int[] b; // array is not static, so skip for now  
  
    public void foo () {  
        a = 0;  
    }  
}
```

▶ C

- global variables and any other variable declared static
- they can be static scalars, arrays or structs or pointers (pointers later)

```
int a;  
int b[10];  
  
void foo () {  
    a = 0;  
    b[a] = a;  
}
```

Static Variable Allocation

```
int a;
int b[10];

void foo () {
    a = 0;
    b[a] = a;
}
```

```
int a;
int b[10];
```

Static Memory Layout

0x1000:	value of a
0x2000:	value of b[0]
0x2004:	value of b[1]
...	
0x2024:	value of b[9]

▶ Key observation

- global/static variables can exist before program starts and live until after it finishes

▶ Static vs dynamic computation

- **compiler** allocates variables, giving them a constant address
- no dynamic computation required to allocate the variables, they just exist
- compiler tracks free space during compilation ... it is in complete control of program's memory

Static Variable Access (scalars)

```
int a;  
int b[10];  
  
void foo () {  
    a = 0;  
    b[a] = a;  
}
```

a = 0;

b[a] = a;

Static Memory Layout

0x1000: value of a
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

▶ Key Observation

- address of **a**, **b[0]**, **b[1]**, **b[2]**, ... are constants known to the compiler

▶ Use RTL to specify instructions needed for **a = 0**

Now Generalize ... What about, for example

a = x + y?
c = a?

Question 1b.1

```
int a;  
int b[10];  
  
void foo () {  
    a = 0;  
    b[a] = a;  
}
```

a = 0;

b[a] = a;

Static Memory Layout

0x1000: value of a
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

- ▶ When is space for **a** allocated (when is its address determined)?
 - A. The compiler assigns the address when it compiles the program
 - B. The compiler calls the memory to allocate **a** when it compiles the program
 - C. The compiler generates code to allocate **a** before the program starts running
 - D. The program locates available space for **a** before it starts running
 - E. The program locates available space for **a** when it starts running
 - F. The program locates available space for **a** just before calling **foo()**

Static Variable Access (static arrays)

```
int a;
int b[10];

void foo () {
    a = 0;
    b[a] = a;
}
```

a = 0;

b[a] = a;

Static Memory Layout

0x1000: value of a
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

▶ Key Observation

- compiler does not know address of **b[a]**
 - unless it can know the value of **a** statically, which it could here by looking at **a=0**, but not in general

▶ Array access is computed from base and index

- address of element is *base* plus *offset*; *offset* is *index* times element size
- the base address (0x2000) and element size (4) are static, the index is dynamic

▶ Use RTL to specify instructions for **b[a] = a**

- assume that the compiler does not know the value of **a** (it does in this case, but not in general)

Choosing the instructions

a = 0;

m[0x1000] ← 0x0

b[a] = a;

m[0x2000 + m[0x1000] × 4] ← m[0x1000]

ISA Design Goals

- ▶ minimize # of memory instructions in ISA
- ▶ at most 1 memory access per instruction
- ▶ minimize # of total instructions in ISA
- ▶ minimize instruction size

r[0] ← 0x0 ①
r[1] ← 0x1000
m[r[1]] ← r[0] ②

r[0] ← 0x1000
r[1] ← m[r[0]] ③
r[2] ← 0x2000
m[r[2] + r[1] × 4] ← r[1] ④

and load version of #4 ⑤

The First Five Instructions

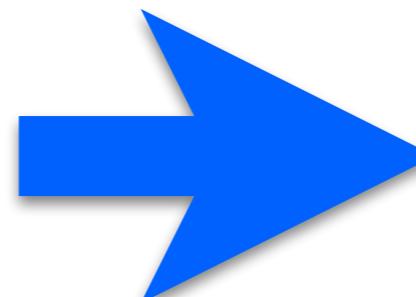
	Name	Semantics	Assembly	Machine
1	<i>load immediate</i>	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
3	<i>load base+offset</i>	$r[d] \leftarrow m[r[s]]$	ld (rs), rd	1?sd
5	<i>load indexed</i>	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd	2sid
2	<i>store base+offset</i>	$m[r[d]] \leftarrow r[s]$	st rs, (rd)	3s?d
4	<i>store indexed</i>	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

Translating the Code

► To RTL

```
int a;
int b[10];

void foo () {
    a = 0;
    b[a] = a;
}
```

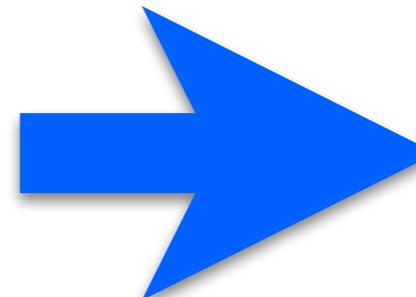


r[0]	← 0
r[1]	← 0x1000
m[r[1]]	← r[0]
r[2]	← m[r[1]]
r[3]	← 0x2000
m[r[3]+r[2]*4]	← r[2]

► To Assembly Code

```
int a;
int b[10];

void foo () {
    a = 0;
    b[a] = a;
}
```



```
ld $0, r0
ld $0x1000, r1
st r0, (r1)

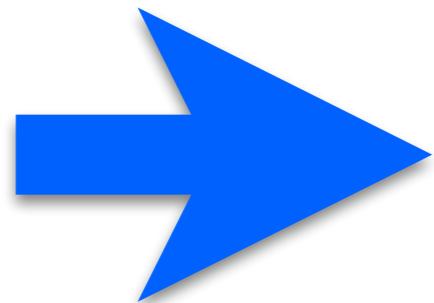
ld (r1), r2
ld $0x2000, r3
st r2, (r3,r2,4)
```

► If a human wrote this assembly

- list static allocations, use labels for addresses, add comments

```
int a;
int b[10];

void foo () {
    a = 0;
    b[a] = a;
}
```



```
ld $0, r0          # r0 = 0
ld $a_data, r1    # r1 = address of a
st r0, (r1)        # a = 0

ld (r1), r2        # r2 = a
ld $b_data, r3    # r3 = address of b
st r2, (r3,r2,4)  # b[a] = a

.pos 0x1000
a_data:
.long 0            # the variable a

.pos 0x2000
b_data:
.long 0            # the variable b[0]
.long 0            # the variable b[1]
...
.long 0            # the variable b[9]
```

The SM213 ISA

(Part 1)

The Simple Machine (SM213) ISA

▶ Architecture

- Register File 8, 32-bit general purpose registers
- CPU one cycle per instruction (fetch + execute)
- Main Memory byte addressed, Big Endian integers

▶ Instruction Format

- 2 or 6 byte instructions (each character is a hex digit)
 - **x-01**, **xxsd**, **x0vv** or **x-sd vvvvvvvv**
- where
 - **x** is an *opcode* (unique identifier for this instruction)
 - **-** means unused
 - **s** and **d** are operand register numbers
 - **vv vvvvvvvv** are immediate / constant values

Machine and Assembly Syntax

▶ Machine code

- [addr:] x-01 [vvvvvvvv]
 - **addr:** sets starting address for subsequent instructions
 - **x-01** hex value of instruction with opcode x and operands 0 and 1
 - **vvvvvvvv** hex value of optional extended value part instruction

▶ Assembly code

- ([label:] [instruction | directive] [# comment] |)*
 - **directive** :: (.pos number) | (.long number)
 - **instruction** :: opcode operand+
 - **operand** :: \$literal | reg | offset (reg) | (reg,reg,4)
 - **reg** :: r 0..7
 - **literal** :: number
 - **offset** :: number
 - **number** :: decimal | 0x hex

Memory-Access Instructions

▶ Load and Store with different *Addressing Modes*

Name	Semantics	Assembly	Machine
<i>load immediate</i>	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
<i>load base+offset</i>	$r[d] \leftarrow m[(o=4*p)+r[s]]$	ld o(rs), rd	1psd
<i>load indexed</i>	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs,ri,4), rd	2sid
<i>store base+offset</i>	$m[(o=4*p)+r[d]] \leftarrow r[s]$	st rs, o(rd)	3spd
<i>store indexed</i>	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd,ri,4)	4sdi

▶ We have specified 4 *addressing modes* for operands

- *immediate* constant value stored in instruction
- *register* operand is register number; register stores value
- *base+offset* operand is register number;
register stores memory address of value (+ offset = p*4)
- *indexed* two register-number operands;
store base memory address and index of value

Basic Arithmetic, Shifting NOP and Halt

► Arithmetic

Name	Semantics	Assembly	Machine
register move	$r[d] \leftarrow r[s]$	mov rs, rd	60sd
add	$r[d] \leftarrow r[d] + r[s]$	add rs, rd	61sd
and	$r[d] \leftarrow r[d] \& r[s]$	and rs, rd	62sd
inc	$r[d] \leftarrow r[d] + 1$	inc rd	63-d
inc address	$r[d] \leftarrow r[d] + 4$	inca rd	64-d
dec	$r[d] \leftarrow r[d] - 1$	dec rd	65-d
dec address	$r[d] \leftarrow r[d] - 4$	deca rd	66-d
not	$r[d] \leftarrow \sim r[d]$	not rd	67-d

► Logical (Shifting), NOP and Halt

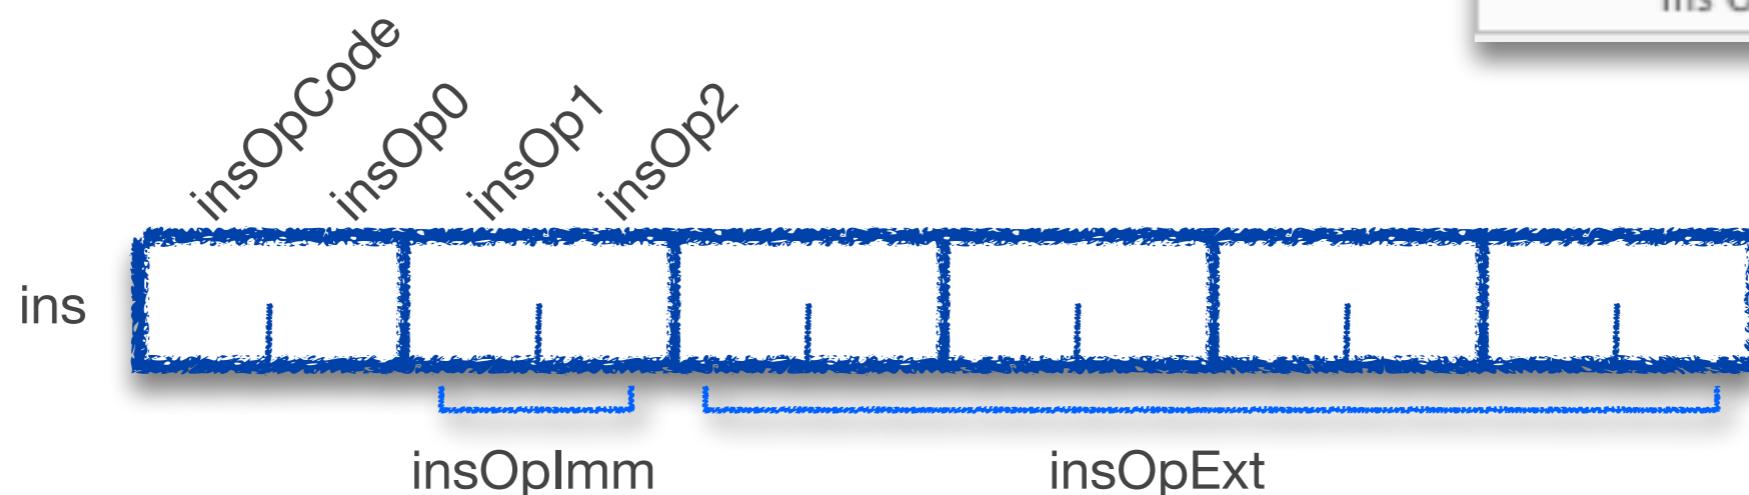
Name	Semantics	Assembly	Machine
shift left	$r[d] \leftarrow r[d] << v=s$	shl \$v, rd	7dss ss>0
shift right	$r[d] \leftarrow r[d] >> v=-s$	shr \$v, rd	7dss ss<0
halt	halt machine	halt	F0--
nop	do nothing	nop	FF--

The CPU Implementation

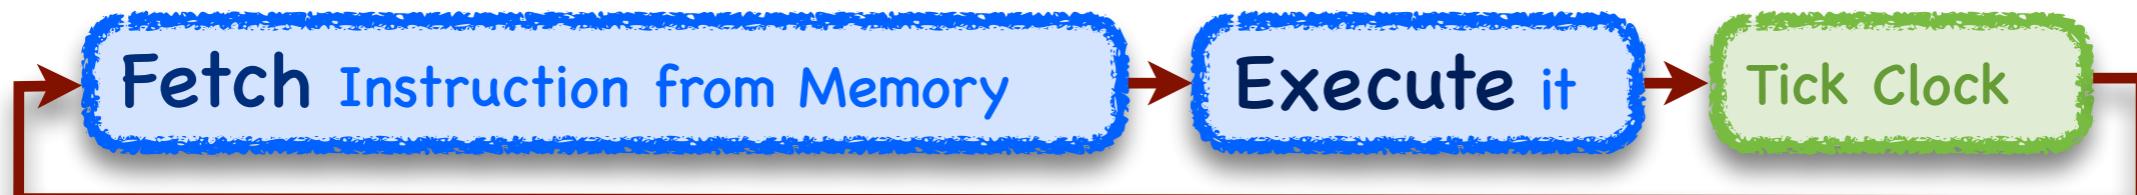
▶ Internal state

- pc address of **next** instruction to fetch
- instruction the value of the current instruction

Reg	Value
PC:	0000010e
Instruction:	3001 00000000
Ins Op Code:	3
Ins Op 0:	0
Ins Op 1:	0
Ins Op 2:	1
Ins Op Imm:	01
Ins Op Ext:	00000000



▶ Cycle stages



- Fetch
 - read instruction at pc, determine size, separate components, set next sequential pc
- Execute
 - read internal state, perform specified computation (insOpCode), update internal state
 - read and update may be to memory as well

Java Syntax

▶ Internal Registers

- insOp0, insOp1, insOp2, insOpImm, insOpExt, pc
- .get()
- .set (value)

```
int i = insOp1.get();
insOp1.set (i);
```

▶ General Purpose Registers

- reg.get (registerNumber)
- reg.set (registerNumber, value)

```
int i = reg.get (3);      // i <= r[3]
reg.set (3, i);          // r[3] <= i
```

▶ Main Memory

- mem.readInteger (address)
- mem.writeInteger (address, value)

```
int i = mem.readInteger (0x1000);      // i <= m[0x1000]
mem.writeInteger (0x1000, i);           // m[0x1000] <= i
```

CPSC 213

Introduction to Computer Systems

Summer Session 2019, Term 2

Unit 1b – July 4, 9

Pointers and Dynamic Arrays

Types

- ▶ In Java & C
 - Every **variable** has a **type**
 - Compiler enforces type safety: restricts what you can do with a variable
- ▶ In Java
 - Types are arranged in a hierarchy (except primitives)
 - Strong type safety: references can't be casted to incompatible types
 - Runtime type information: every object carries its type in memory
- ▶ In C
 - A type defines the *size* and *interpretation* of a bunch of bytes
 - `sizeof` operator can determine the size *at compile-time*
 - No runtime type information: all that matters is how code accesses bytes
- ▶ In Assembly
 - No types! Everything is a number! Instructions decide how to interpret them.

Dynamic Arrays

Global Dynamic Array

- ▶ **Java:** array variable stores reference to array allocated dynamically with **new** statement

```
public class Foo {  
    static int a;  
    static int b[];  
  
    void foo() {  
        b = new int[10];  
        b[a] = a;  
    }  
}
```

- ▶ **C:** array variables can store static arrays or pointers to arrays allocated dynamically with call to **malloc** library procedure

```
int a;  
int *b; # of bytes to allocate  
  
void foo() {  
    b = malloc(10 * sizeof(int));  
    b[a] = a;  
}
```

Static vs Dynamic Arrays

- ▶ Declared and allocated differently, but accessed the same

```
int a;  
int b[10];  
  
void foo() {  
    b[a] = a;  
}
```

```
int a;  
int *b;  
  
void foo() {  
    b = malloc(10 * sizeof(int));  
    b[a] = a;  
}
```

- ▶ Static allocation

- for static arrays, the compiler allocates the array
- for dynamic arrays, the compiler allocates a pointer

0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

0x2000: value of b

▶ Then when the program runs

- the dynamic array is allocated by a call to malloc, say at address 0x3000
- the value of variable b is set to the memory address of this array

0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

0x2000: 0x3000
0x3000: value of b[0]
0x3004: value of b[1]
...
0x3024: value of b[9]



▶ Generating code to access the array

- for the dynamic array, the compiler generates an additional load for b

Static Array

```
r[0] ← 0x1000  
r[1] ← m[r[0]]  
r[2] ← 0x2000  
m[r[2]+r[1]*4] ← r[1]
```

Dynamic Array

```
r[0] ← 0x1000  
r[1] ← m[r[0]]  
r[2] ← 0x2000  
r[3] ← m[r[2]]  
m[r[3]+r[1]*4] ← r[1]
```

load a
load b
b[a]=a

malloc

- ▶ Dumb but functional, illustrative implementation

```
int next = 0;
int memory[100000000];

void *malloc(int size) {
    int *result = &memory[next];
    next += (size + 3) / 4;
    return result;
}
```

- ▶ The hard part is *free* - releasing memory you don't need

How C Arrays are Different from Java

▶ Terminology

- use the term **pointer** instead of **reference**; they mean the same thing

▶ Declaration

- the type is a pointer to the type of its elements, indicated with a *

▶ Allocation

- malloc allocates a block of bytes; no type; no constructor

▶ Type Safety

- any pointer can be type cast to any pointer type

▶ Bounds checking

- C performs no array bounds checking
- out-of-bounds access manipulates memory that is not part of array
- this is the major source of virus vulnerabilities in the world today

Question: Can array bounds checking be performed statically?

* what does this say about a tradeoff that Java and C take differently?

```
void izero(int *a, int n) {
    for(int i=0; i<n; i++)
        a[i] = 0;
}
```

Array Access in Assembly

Static Array

```
ld $a, r0          # r0 = address of a
ld (r0), r1        # r1 = a
ld $b, r2          # r2 = address of b
st r1, (r2,r1,4)  # b[a] = a

.pos 0x1000
a:
.long 0           # the variable a

.pos 0x2000
b:
.long 0           # the variable b[0]
.long 0           # the variable b[1]
...
.long 0           # the variable b[9]
```

Dynamic Array

```
ld $a, r0          # r0 = address of a
ld (r0), r1        # r1 = a
ld $b, r2          # r2 = address of b
ld (r2), r3      # r3 = b == &b[0]
st r1, (r3,r1,4)  # b[a] = a

.pos 0x1000
a:
.long 0           # the variable a

.pos 0x2000
b:
.long 0           # the variable b
```

For Simulator ...

For the simulator, we give the allocation of `b_data` even though this would have been allocated dynamically, and give `b` a value even though this would have happened while the program was running. And so, while the compiler would have assigned `b` the value 0 as you see above, in the snippet you'll see "`b: .long b_data`" instead.

```
.pos 0x3000
b_data:
.long 0           # the variable b[0]
.long 0           # the variable b[1]
...
.long 0           # the variable b[9]
```

Pointers

- ▶ In both languages
 - A pointer/reference variable is **just a number** - a memory address
 - A pointer/reference type refers to another type
- ▶ In Java
 - References can be casted up and down the hierarchy (Object is the root)
 - Runtime checks type during casting
 - References can't be used like numbers, compiler forbids it
- ▶ In C
 - Pointers can be casted to other pointers or even numbers
 - Compiler lets you treat pointers like numbers
 - No checks in casting at all (none from compiler or runtime)

Pointers

- ▶ Pointers are just numbers (addresses)!
- ▶ Pointer type determines size & meaning

C and Java Arrays and Pointers

▶ In both languages

- an array is a list of items of the same type
- array elements are named by non-negative integers starting with 0
- syntax for accessing element **i** of array **b** is **b[i]**

▶ In Java

- variable **a** stores a pointer to the array
- **b[x] = 0** means $m[m[b] + x * \text{sizeof(array-element)}] \leftarrow 0$

▶ In C

- variable **a** can store a pointer to the array or the array itself
- **b[x] = 0** means $m[b + x * \text{sizeof(array-element)}] \leftarrow 0$
or $m[m[b] + x * \text{sizeof(array-element)}] \leftarrow 0$
- dynamic arrays are just like all other pointers
 - stored in **TYPE***
 - **access with either a[x] or *(a+x)**

Example

- The following C programs are identical

```
int *a;  
a[4] = 5;
```

```
int *a;  
*(a+4) = 5;
```

- For array access, the compiler would generate this code

r[0]	← a
r[1]	← m[r[0]]
r[2]	← 4
r[3]	← 5
m[r[1]+r[2]*4]	← r[3]

```
ld $a, r0  
ld (r0), r1  
ld $4, r2  
ld $5, r3  
st r3, (r1,r2,4)
```

- multiplying the index 4 by 4 (size of integer) to compute the array offset
- So, what does this tell you about pointer arithmetic in C?

Adding X to a pointer of type Y*, adds X * sizeof(Y)
to the pointer's memory-address value.

Pointer Arithmetic in C

▶ Its purpose

- an alternative way to access dynamic arrays compared to `a[i]`

▶ Adding or subtracting an integer index to a pointer

- results in a new pointer of the same type
- value of the pointer is offset by index times size of pointer's referent
- for example
 - adding 3 to an `int*` yields a pointer value 12 larger than the original

▶ Subtracting two pointers of the same type

- results in an integer
- gives number of referent-type elements between the two pointers
- for example
 - `(& a[7]) - (& a[2]) == 5 == (a+7) - (a+2)`

▶ other operators

- `& x` the address of X
 - `&a[0]` (or just `a` with no `&` and no `[]`) equals the address of the first element of the array
- `* x` the value X points to

Question 1b.2

```
int *c;

void foo () {
    // ...
    c = (int *) malloc (10*sizeof(int));
    // ...
    c = &c[3];
    *c = *&c[3];
    // ...
}
```

- ▶ What is the equivalent Java statement to
 - *that is, which one makes the same change to the array C*
 - c[0] = c[3];
 - c[3] = c[6];
 - there is no type-safe equivalent
 - not valid, because you can't take the address of a static in Java

Looking more closely

```
c = &c[3];  
*c = *&c[3];
```

r[0] ← 0x2000	# r[0] = &c
r[1] ← m[r[0]]	# r[1] = c
r[2] ← 12	# r[2] = 3 * sizeof(int)
r[2] ← r[2]+r[1]	# r[2] = c + 3
m[r[0]] ← r[2]	# c = c + 3
r[3] ← 3	# r[3] = 3
r[4] ← m[r[2]+4*r[3]]	# r[4] = c[3]
m[r[2]] ← r[4]	# c[0] = c[3]

Before

0x2000: 0x3000			<table border="1"><tbody><tr><td>0x3000: 0</td></tr><tr><td>0x3004: 1</td></tr><tr><td>0x3008: 2</td></tr><tr><td>0x300c: 3</td></tr><tr><td>0x3010: 4</td></tr><tr><td>0x3014: 5</td></tr><tr><td>0x3018: 6</td></tr><tr><td>0x301c: 7</td></tr><tr><td>0x3020: 8</td></tr><tr><td>0x3024: 9</td></tr></tbody></table>	0x3000: 0	0x3004: 1	0x3008: 2	0x300c: 3	0x3010: 4	0x3014: 5	0x3018: 6	0x301c: 7	0x3020: 8	0x3024: 9
0x3000: 0													
0x3004: 1													
0x3008: 2													
0x300c: 3													
0x3010: 4													
0x3014: 5													
0x3018: 6													
0x301c: 7													
0x3020: 8													
0x3024: 9													

After

0x2000: 0x300c			<table border="1"><tbody><tr><td>0x3000: 0</td></tr><tr><td>0x3004: 1</td></tr><tr><td>0x3008: 2</td></tr><tr><td>0x300c: 6</td></tr><tr><td>0x3010: 4</td></tr><tr><td>0x3014: 5</td></tr><tr><td>0x3018: 6</td></tr><tr><td>0x301c: 7</td></tr><tr><td>0x3020: 8</td></tr><tr><td>0x3024: 9</td></tr></tbody></table>	0x3000: 0	0x3004: 1	0x3008: 2	0x300c: 6	0x3010: 4	0x3014: 5	0x3018: 6	0x301c: 7	0x3020: 8	0x3024: 9
0x3000: 0													
0x3004: 1													
0x3008: 2													
0x300c: 6													
0x3010: 4													
0x3014: 5													
0x3018: 6													
0x301c: 7													
0x3020: 8													
0x3024: 9													
 			<table border="1"><tbody><tr><td>0x3000: 0</td></tr><tr><td>0x3004: 1</td></tr><tr><td>0x3008: 2</td></tr><tr><td>0x300c: 6</td></tr><tr><td>0x3010: 4</td></tr><tr><td>0x3014: 5</td></tr><tr><td>0x3018: 6</td></tr><tr><td>0x301c: 7</td></tr><tr><td>0x3020: 8</td></tr><tr><td>0x3024: 9</td></tr></tbody></table>	0x3000: 0	0x3004: 1	0x3008: 2	0x300c: 6	0x3010: 4	0x3014: 5	0x3018: 6	0x301c: 7	0x3020: 8	0x3024: 9
0x3000: 0													
0x3004: 1													
0x3008: 2													
0x300c: 6													
0x3010: 4													
0x3014: 5													
0x3018: 6													
0x301c: 7													
0x3020: 8													
0x3024: 9													

► And in assembly language

r[0] ← 0x2000	# r[0] = &c
r[1] ← m[r[0]]	# r[1] = c
r[2] ← 12	# r[2] = 3 * sizeof(int)
r[2] ← r[2]+r[1]	# r[2] = c + 3
m[r[0]] ← r[2]	# c = c + 3
r[3] ← 3	# r[3] = 3
r[4] ← m[r[2]+4*r[3]]	# r[4] = c[3]
m[r[2]] ← r[4]	# c[0] = c[3]

ld \$0x2000, r0	# r0 = &c
ld (r0), r1	# r1 = c
ld \$12, r2	# r2 = 3*sizeof(int)
add r1, r2	# r2 = c+3
st r2, (r0)	# c = c+3
ld \$3, r3	# r3 = 3
ld (r2,r3,4), r4	# r4 = c[3]
st r4, (r2)	# c[0] = c[3]

Aside: Strings in C

- ▶ A3 shows real C programs: `int main(int argc, char **argv)`
- ▶ `argc` = *count* of arguments
- ▶ `argv` = *vector* (array) of arguments. `argv[0]` = program name
- ▶ `char *` = array of characters (1 byte each) = *string*
 - Terminated by a NUL character
 - e.g. "Hello World!" =

H	e	I	I	o		W	o	r	I	d	!	\0
---	---	---	---	---	--	---	---	---	---	---	---	----
 - Characters translated to numbers using ASCII
- ▶ `char **` = array of strings
- ▶ Use `atoi` or `strtol` to convert string to int

Aside: I/O in C

► *printf*: Print Formatted - print some data using format

- %s = string
- %d = decimal (signed) integer
- %u = unsigned integer
- %x = hexadecimal integer
- For each format specifier, add one function argument
- *fprintf*: print to file (e.g. stderr = standard error, like System.err)
- Example: `printf("My name is %s and I'm %d years old", "Robert", 30);`

► *scanf*: Scan Formatted - read some data using format

- All arguments must be pointers!

Aside: Declaring Pointers in C

- ▶ What happens when you want to declare multiple pointers?
 - To declare multiple ints, just `int a, b;`
 - To declare multiple pointers, what about `int *a, b;?`
- ▶ No! The `*` belongs to each variable in the line!
- ▶ Correct way to declare multiple pointers
 - `int *a, *b;`
- ▶ Can combine multiple levels of pointers and initialization...
 - `int a=42, *b=&a, **c=&b;`
 - Don't do this at home
- ▶ Same thing applies to arrays
 - `int a, b[5], c[10];`

Summary: Static Scalar and Array Variables

- ▶ Static variables
 - the compiler knows the address (memory location) of variable
- ▶ Static scalars and arrays
 - the compiler knows the address of the scalar value or array
- ▶ Dynamic arrays
 - the compiler does not know the address the array
- ▶ What C does that Java doesn't
 - static arrays
 - arrays can be accessed using pointer dereferencing operator
 - arithmetic on pointers
- ▶ What Java does that C doesn't
 - type-safe dynamic allocation
 - automatic array-bounds checking