

CPSC 213

Introduction to Computer Systems

Summer Session 2019, Term 2

Unit 2a - Jul 30

I/O Devices, Interrupts and DMA

Overview


▶ Reading in Text

- 8.1, 8.2.1, 8.5.1-8.5.3

▶ Learning Goals

- Explain what PIO is, why it exists, what processor originates it, and how it is used
- Explain what DMA is, why it exists, what it does, and what processor originates it
- Explain what an interrupt is, why it exists, what it does, and what processor originates it
- Compare PIO and DMA by identifying when each *can* be used and when each *should* be used
- Explain the relationship between polling, PIO and interrupts
- Explain the conditions that make polling acceptable or not
- Explain what happens when an interrupt occurs at the hardware level
- Explain what it means for an operation such as a disk read to be asynchronous and give examples of code that works when an operation is synchronous but does not work with it is asynchronous
- Write event-driven programs in C using function pointers
- Describe why event-driven programs may be harder to write, read, and debug

Looking Beyond the CPU and Memory



▶ Memory Bus

- data/control path connecting CPU, Main Memory, and I/O Bus
- also called the *Front Side Bus*

▶ I/O Bus

- data/control path connecting Memory Bus and I/O Controllers
- e.g., PCI

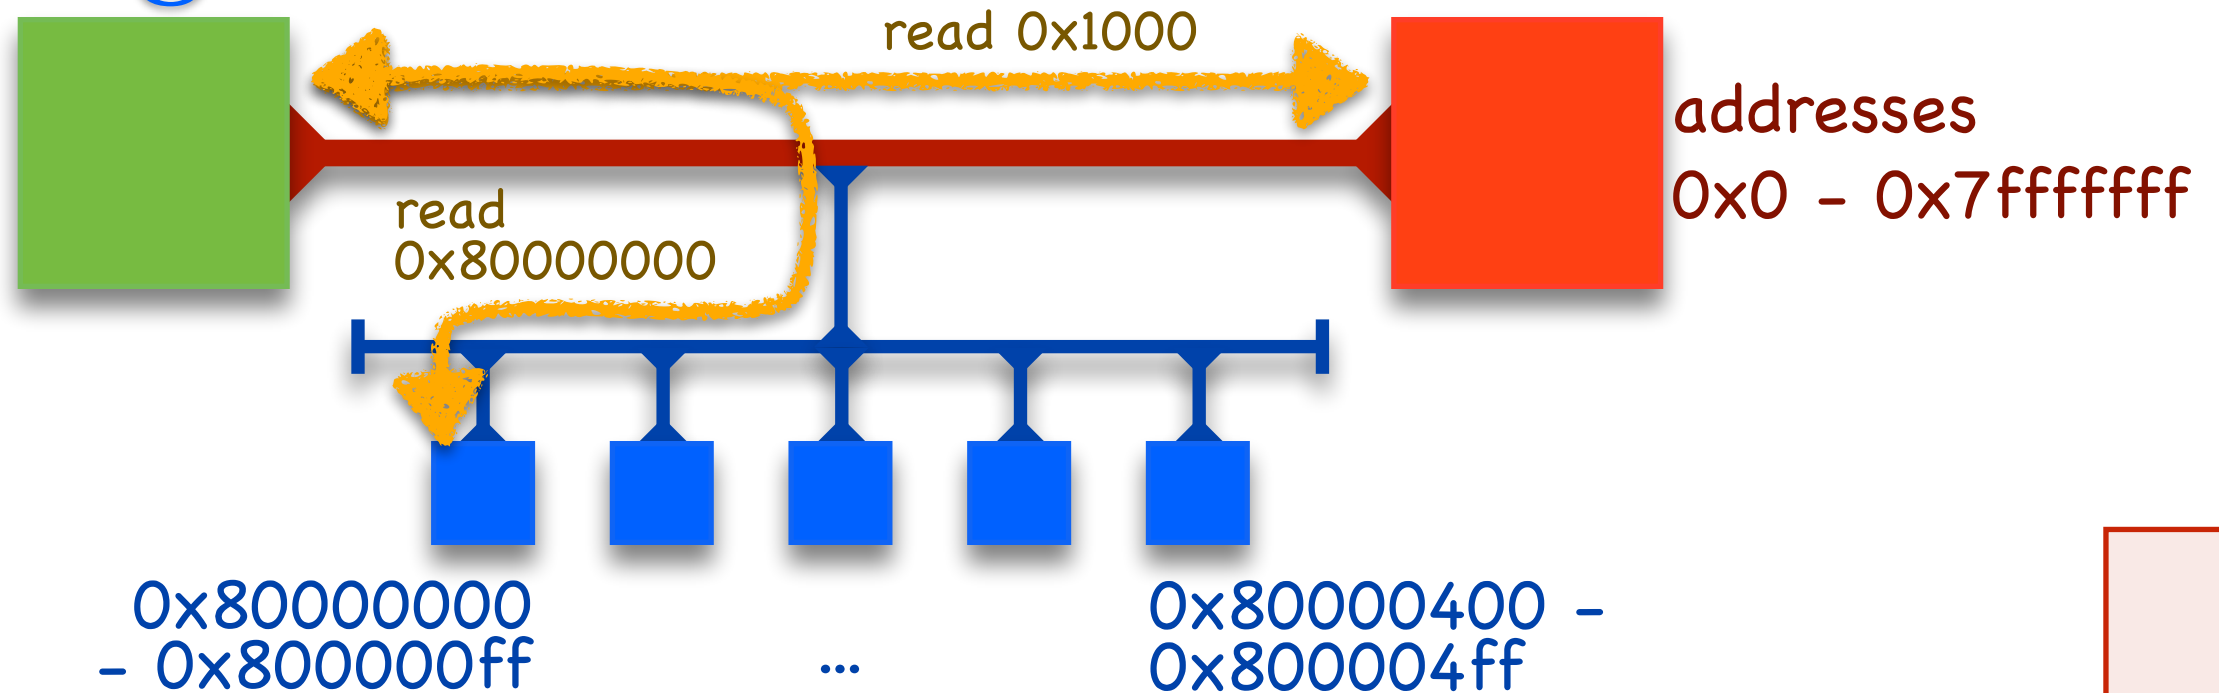
▶ I/O Controller

- a processor running software (firmware)
- connects I/O Device to I/O Bus
- e.g. ,SCSI, SATA, Ethernet, ...

▶ I/O Device

- I/O mechanism that generates or consumes data
- e.g., disk, radio, keyboard, mouse, ...

Talking to an I/O Controller



▶ Programmed I/O (PIO)

- CPU transfers a word at a time between CPU and I/O controller
- typically use standard load/store instructions, but to I/O-mapped memory

▶ I/O-Mapped Memory

- memory addresses outside of main memory
- used to name I/O controllers (usually configured at boot time)
- loads and stores are translated into I/O-bus messages to controller

▶ Example

- to read/write to controller at address 0x80000000

```
ld    $0x80000000, r0
st    r1 (r0)          # write the value of r1 to the device
ld    (r0), r1         # read a word from device into r1
```



Limitations of PIO

- ▶ Reading or writing large amounts of data slows CPU
 - requires CPU to transfer one word at a time
 - controller/device is much slower than CPU
 - and so, CPU runs at controller/device speed, mostly waiting for controller
- ▶ IO Controller cannot initiate communication
 - sometimes the CPU asks for data
 - but, sometimes controller receives data for the CPU, without CPU asking
 - e.g., mouse click or network packet reception (everything is like this really as we will see)
 - how does controller notify CPU that it has data the CPU should want?
- ▶ One idea...
 - what is it? _____
 - what are drawbacks? _____
 - when is it okay? _____

Polling and I/O Memory



```
int pollDeviceForValue() {
    volatile int *ready = 0x800000100;
    volatile int *value = 0x800000104;
    int v;

    while(*ready == 0) {}; (1)
    v = *value; (4)
    *ready = 0;

    return v;
}
```

```
void readyValueForCPUPoll(int v) {
    volatile int *ready = 0x800000100;
    volatile int *value = 0x800000104;

    while(*ready != 0) {};
    *value = v; (2)
    *ready = 1; (3)
}

readyValueForCPUPoll (7);
```


Reading (or writing) I/O Memory IS SLOW

CPU → I/O Device: give me value at address X

I/O Device → CPU: here is value

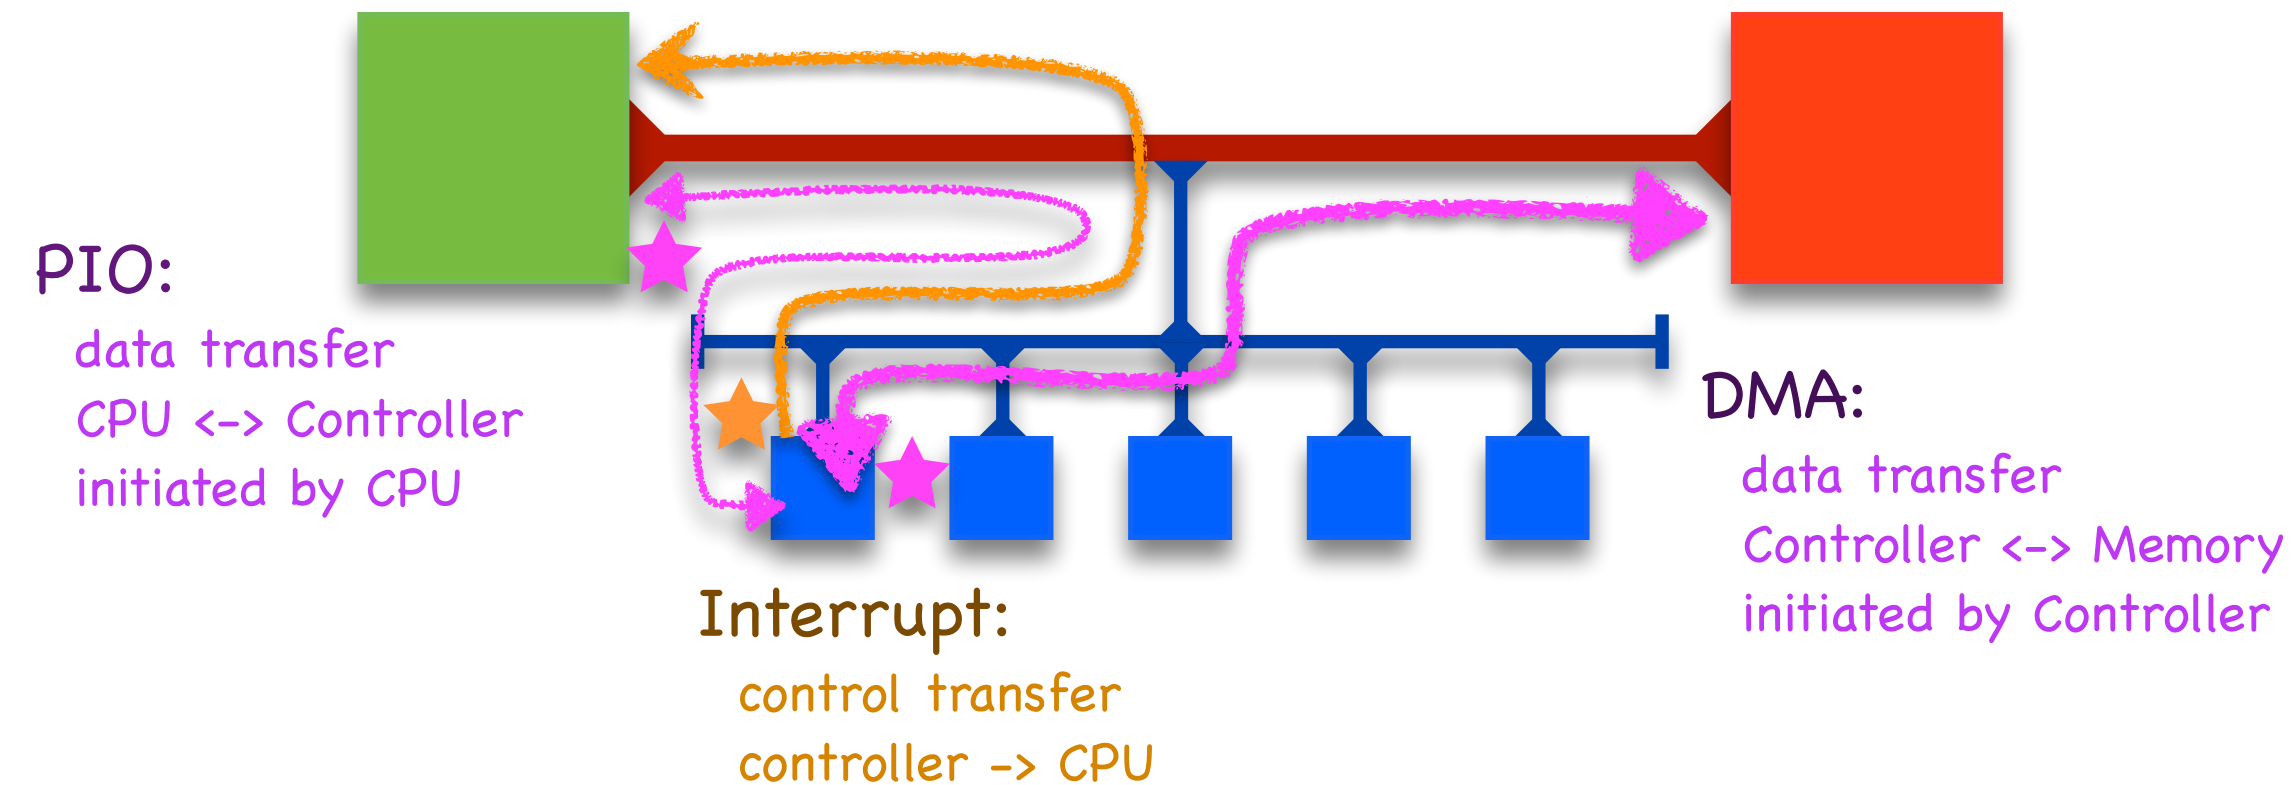
Polling is Okay If poll has low overhead or if high-probability of “hit”.

Key Observation



- ▶ CPU and I/O Controller are independent processors
 - they should be permitted to work in parallel
 - either should be able to initiate data transfer to/from memory
 - either should be able to signal the other to get the other's attention

Autonomous Controller Operation



▶ Direct Memory Access (DMA)

- controller can send/read data from/to any main memory address
- the CPU is oblivious to these transfers
- DMA addresses and sizes are *programmed* by CPU using PIO

▶ CPU Interrupts

- controller can signal the CPU
- CPU checks for interrupts on every cycle (its like a really fast, clock-speed poll)
- CPU jumps to controller's *Interrupt Service Routine* if it is interrupting

Adding Interrupts to Simple CPU

▶ New special-purpose CPU registers

- **isDeviceInterrupting** set by I/O Controller to signal interrupt
- **interruptControllerID** set by I/O Controller to identify interrupting device
- **interruptVectorBase** interrupt-handler jump table, initialized a boot time

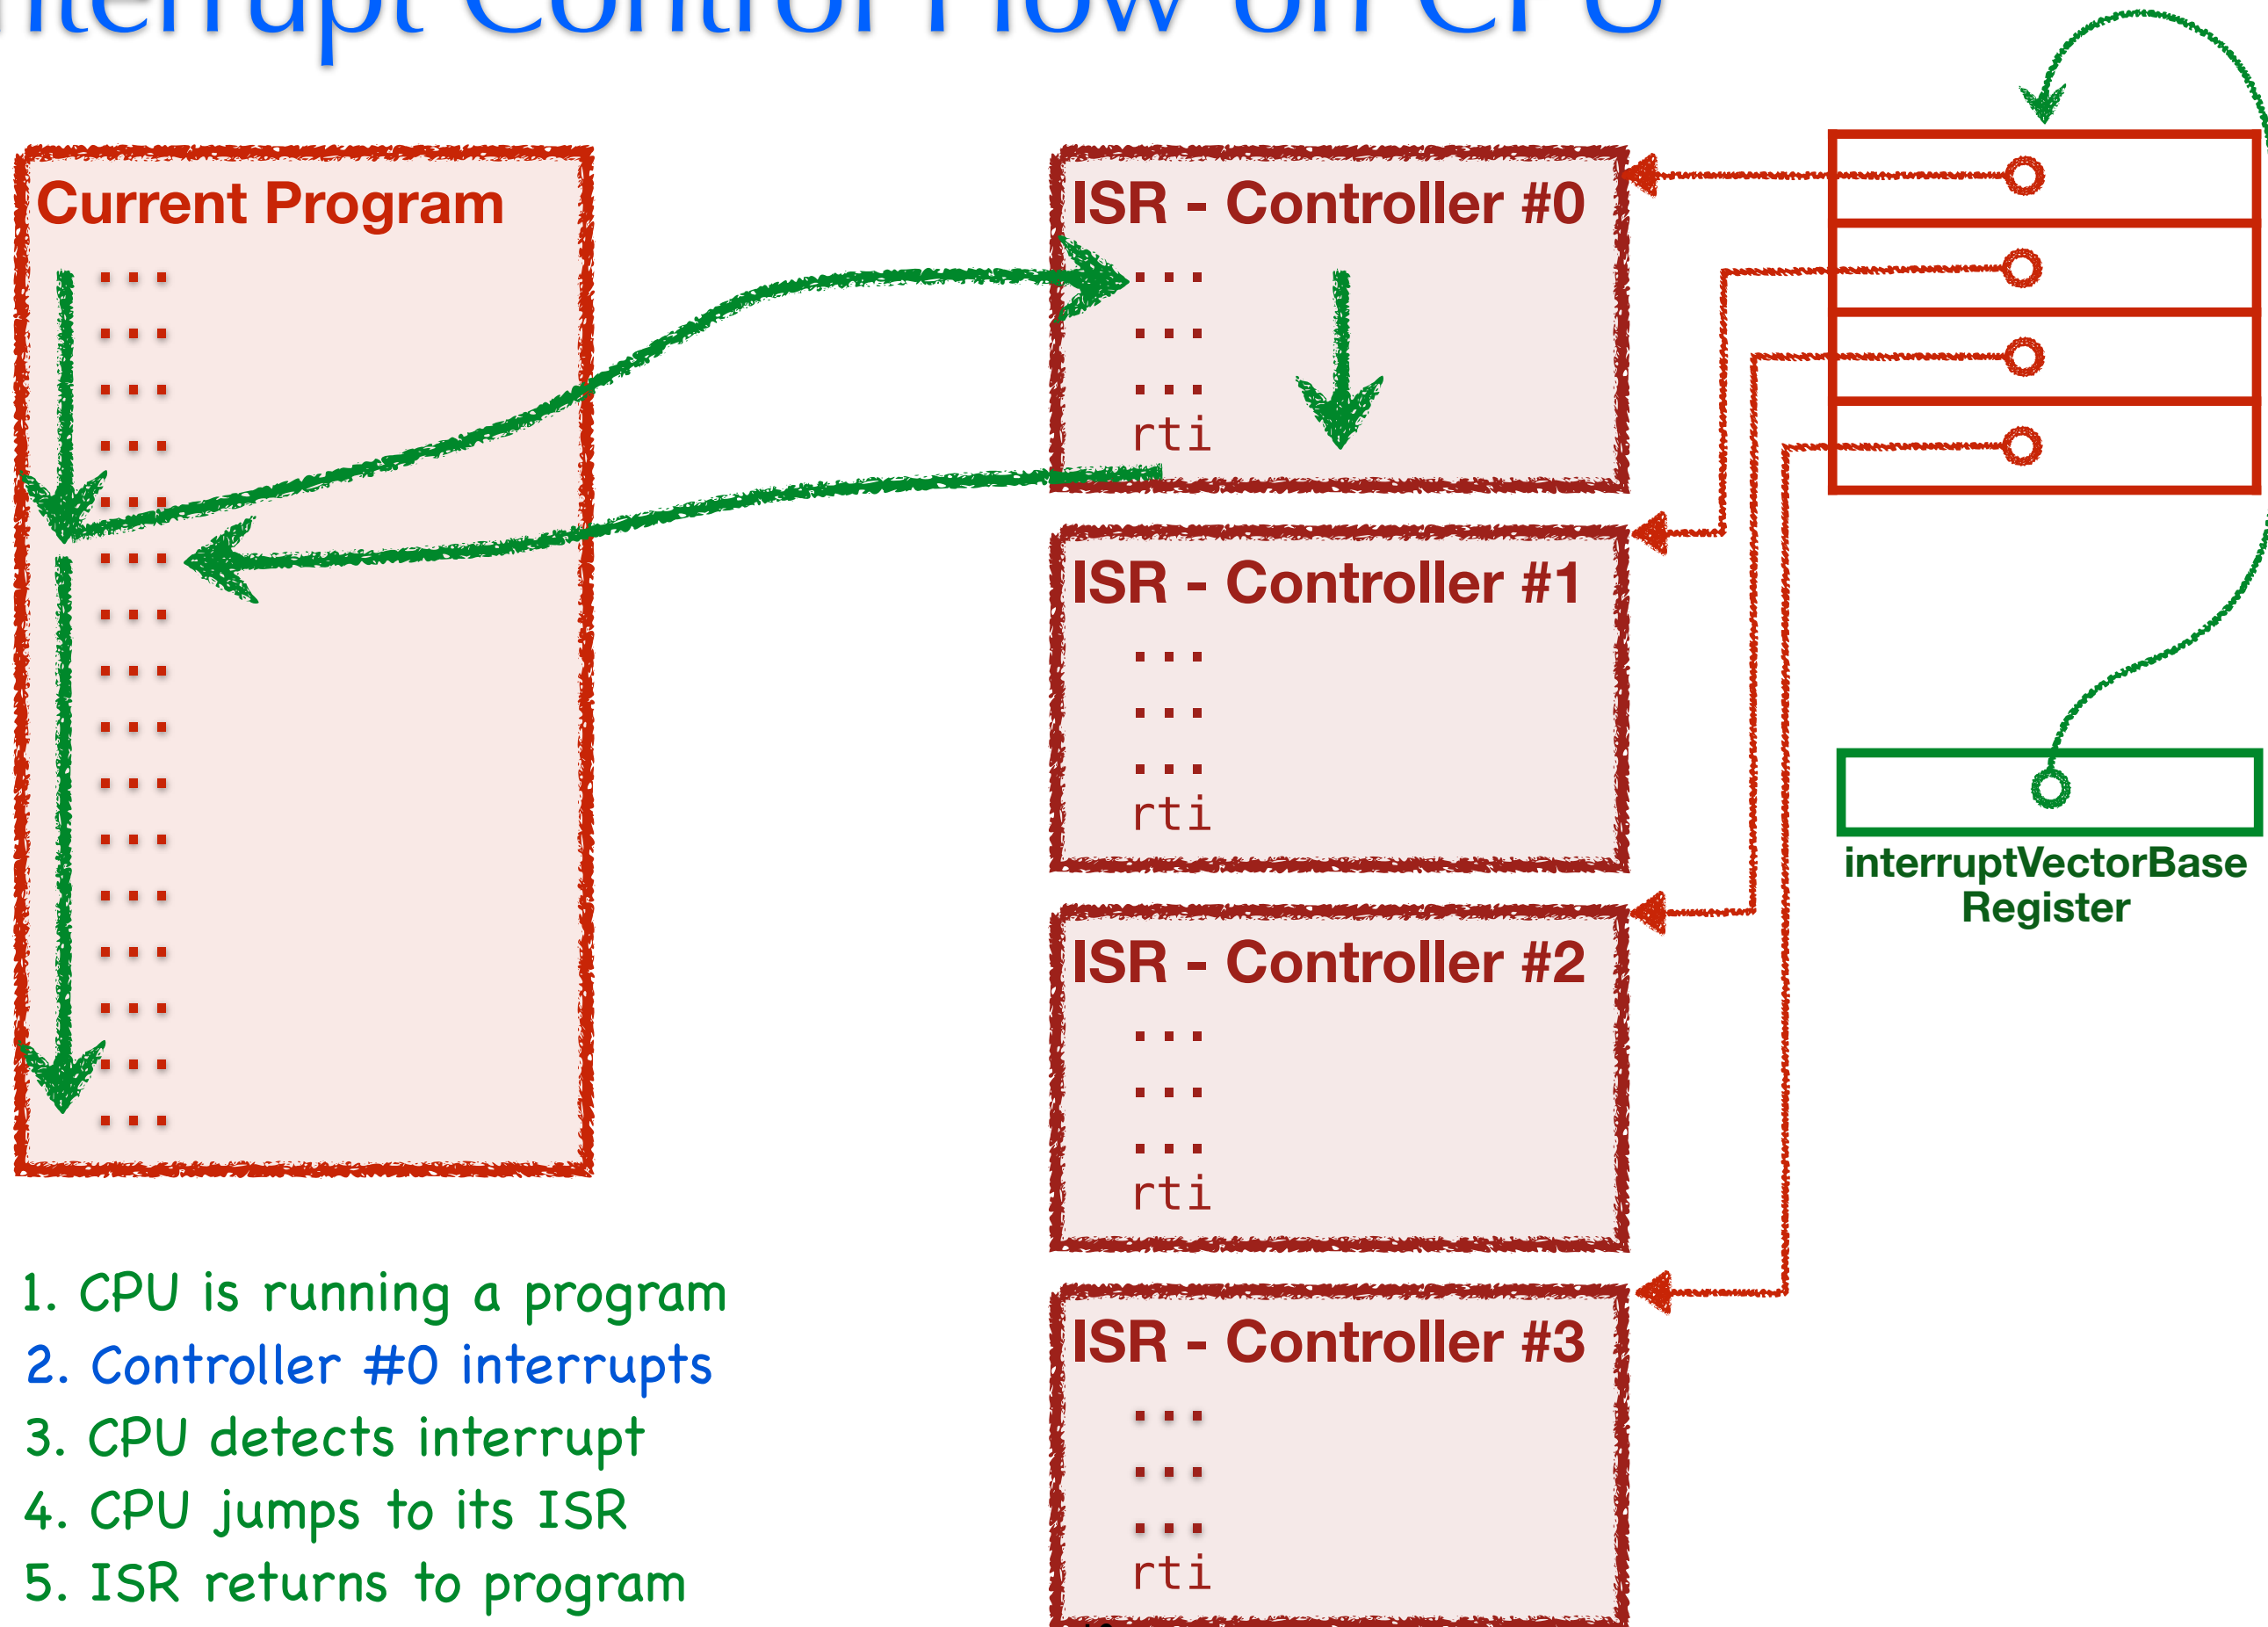
▶ Modified fetch-execute cycle

```
while(true) {  
    if(isDeviceInterrupting) {  
        r[5] ← r[5] - 4; m[r[5]] ← r[6];  
        r[6] ← pc;  
        pc ← interruptVectorBase[interruptControllerID];  
    }  
    fetch();  
    execute();  
}
```


RTI: Return from Interrupt Instruction

```
t ← r[6];  
r[6] ← m[r[5]]; r[5] ← r[5] + 4;  
isDeviceInterrupting ← 0;  
pc ← t;
```

Interrupt Control Flow on CPU



Architectural View of Interrupts



```
while(true) {  
    if(isDeviceInterrupting) {  
        r[5] ← r[5] - 4; m[r[5]] ← r[6];  
        r[6] ← pc;  
        pc ← interruptVectorBase[interruptControllerID];  
        isDeviceInterrupting ← 0;  
    }  
    fetch();  
    execute();  
}
```

Polling on CPU register instead of I/O Memory: turning bad polling into good polling

Programming with I/O

Disk Read Timeline

CPU

1. PIO to request read

...

do other things

...

6. Interrupt Received
Call readComplete

I/O Controller

2. PIO Received, start read

...

wait for read to complete

...

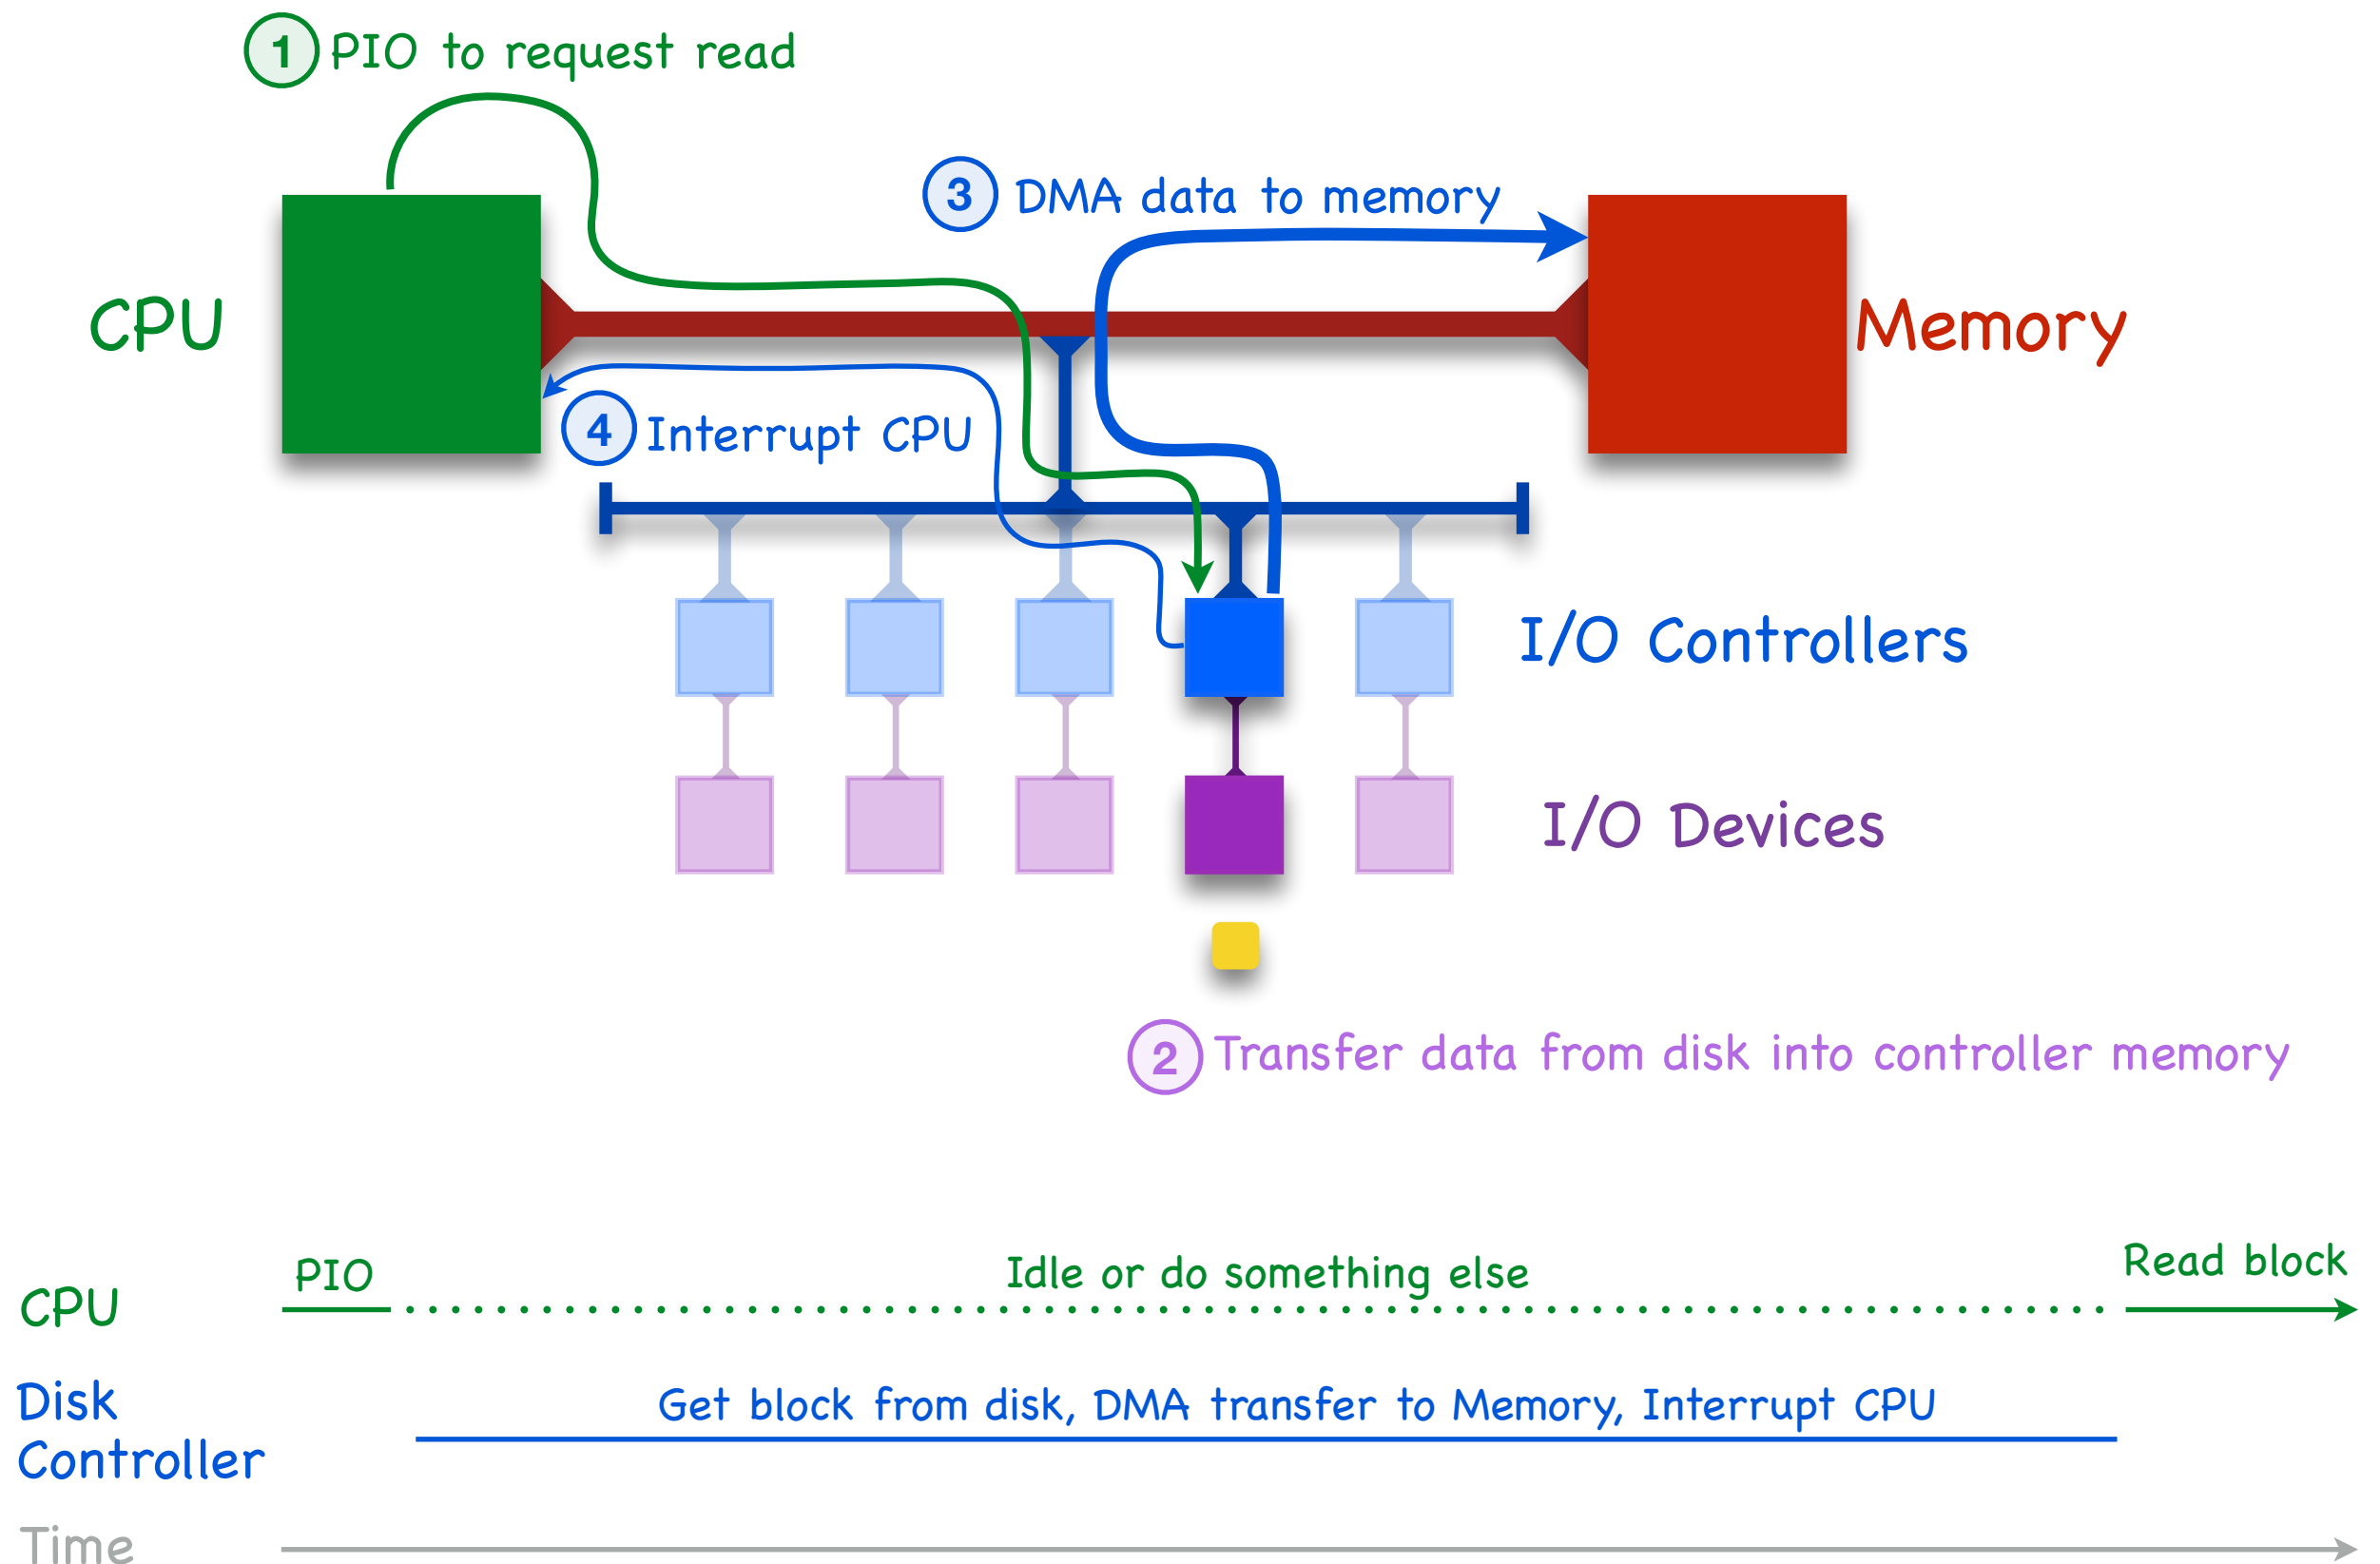
3. Read completes

4. Transfer data to memory (DMA)

5. Interrupt CPU

A disk stores blocks of data. Each block has a number. CPU can read or write.

Disk Read Timeline



The Power of the Sequence

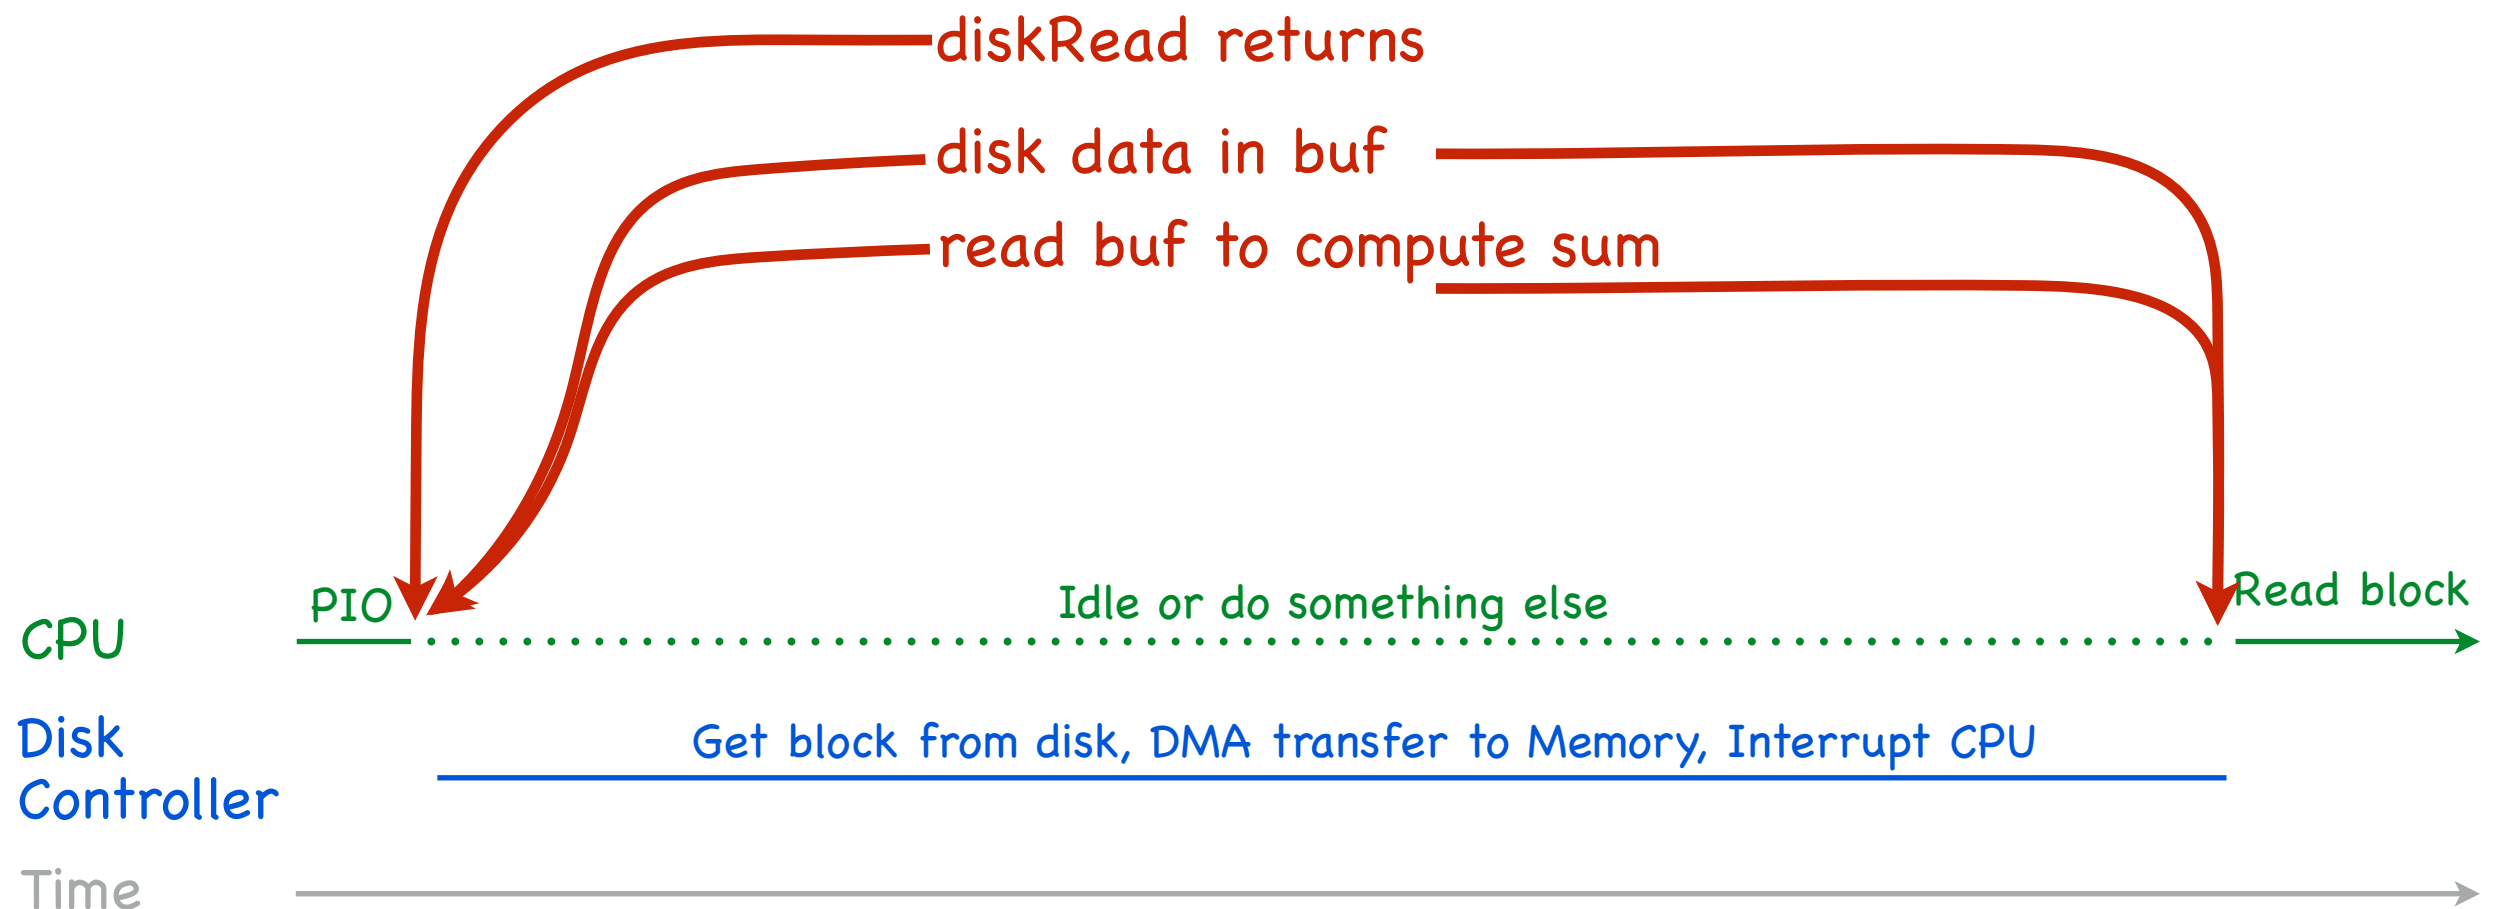
- ▶ Consider a program that reads data from a disk
 - you can sort of think of the read has having three parts:
 1. figure out what data you want
 2. get the data
 3. do something with the data you got
 - these steps must happen in this order
 - we think of these steps as a sequence
- ▶ For Example ...

```
int sumDiskData(blockNum, numBytes) {  
    char buf[numBytes];  
    int sum = 0;  
  
    diskRead(buf, blockNum, numBytes);  
    for(int i=0; i<numBytes; i++)  
        sum += buf[i];  
  
    return sum;  
}
```

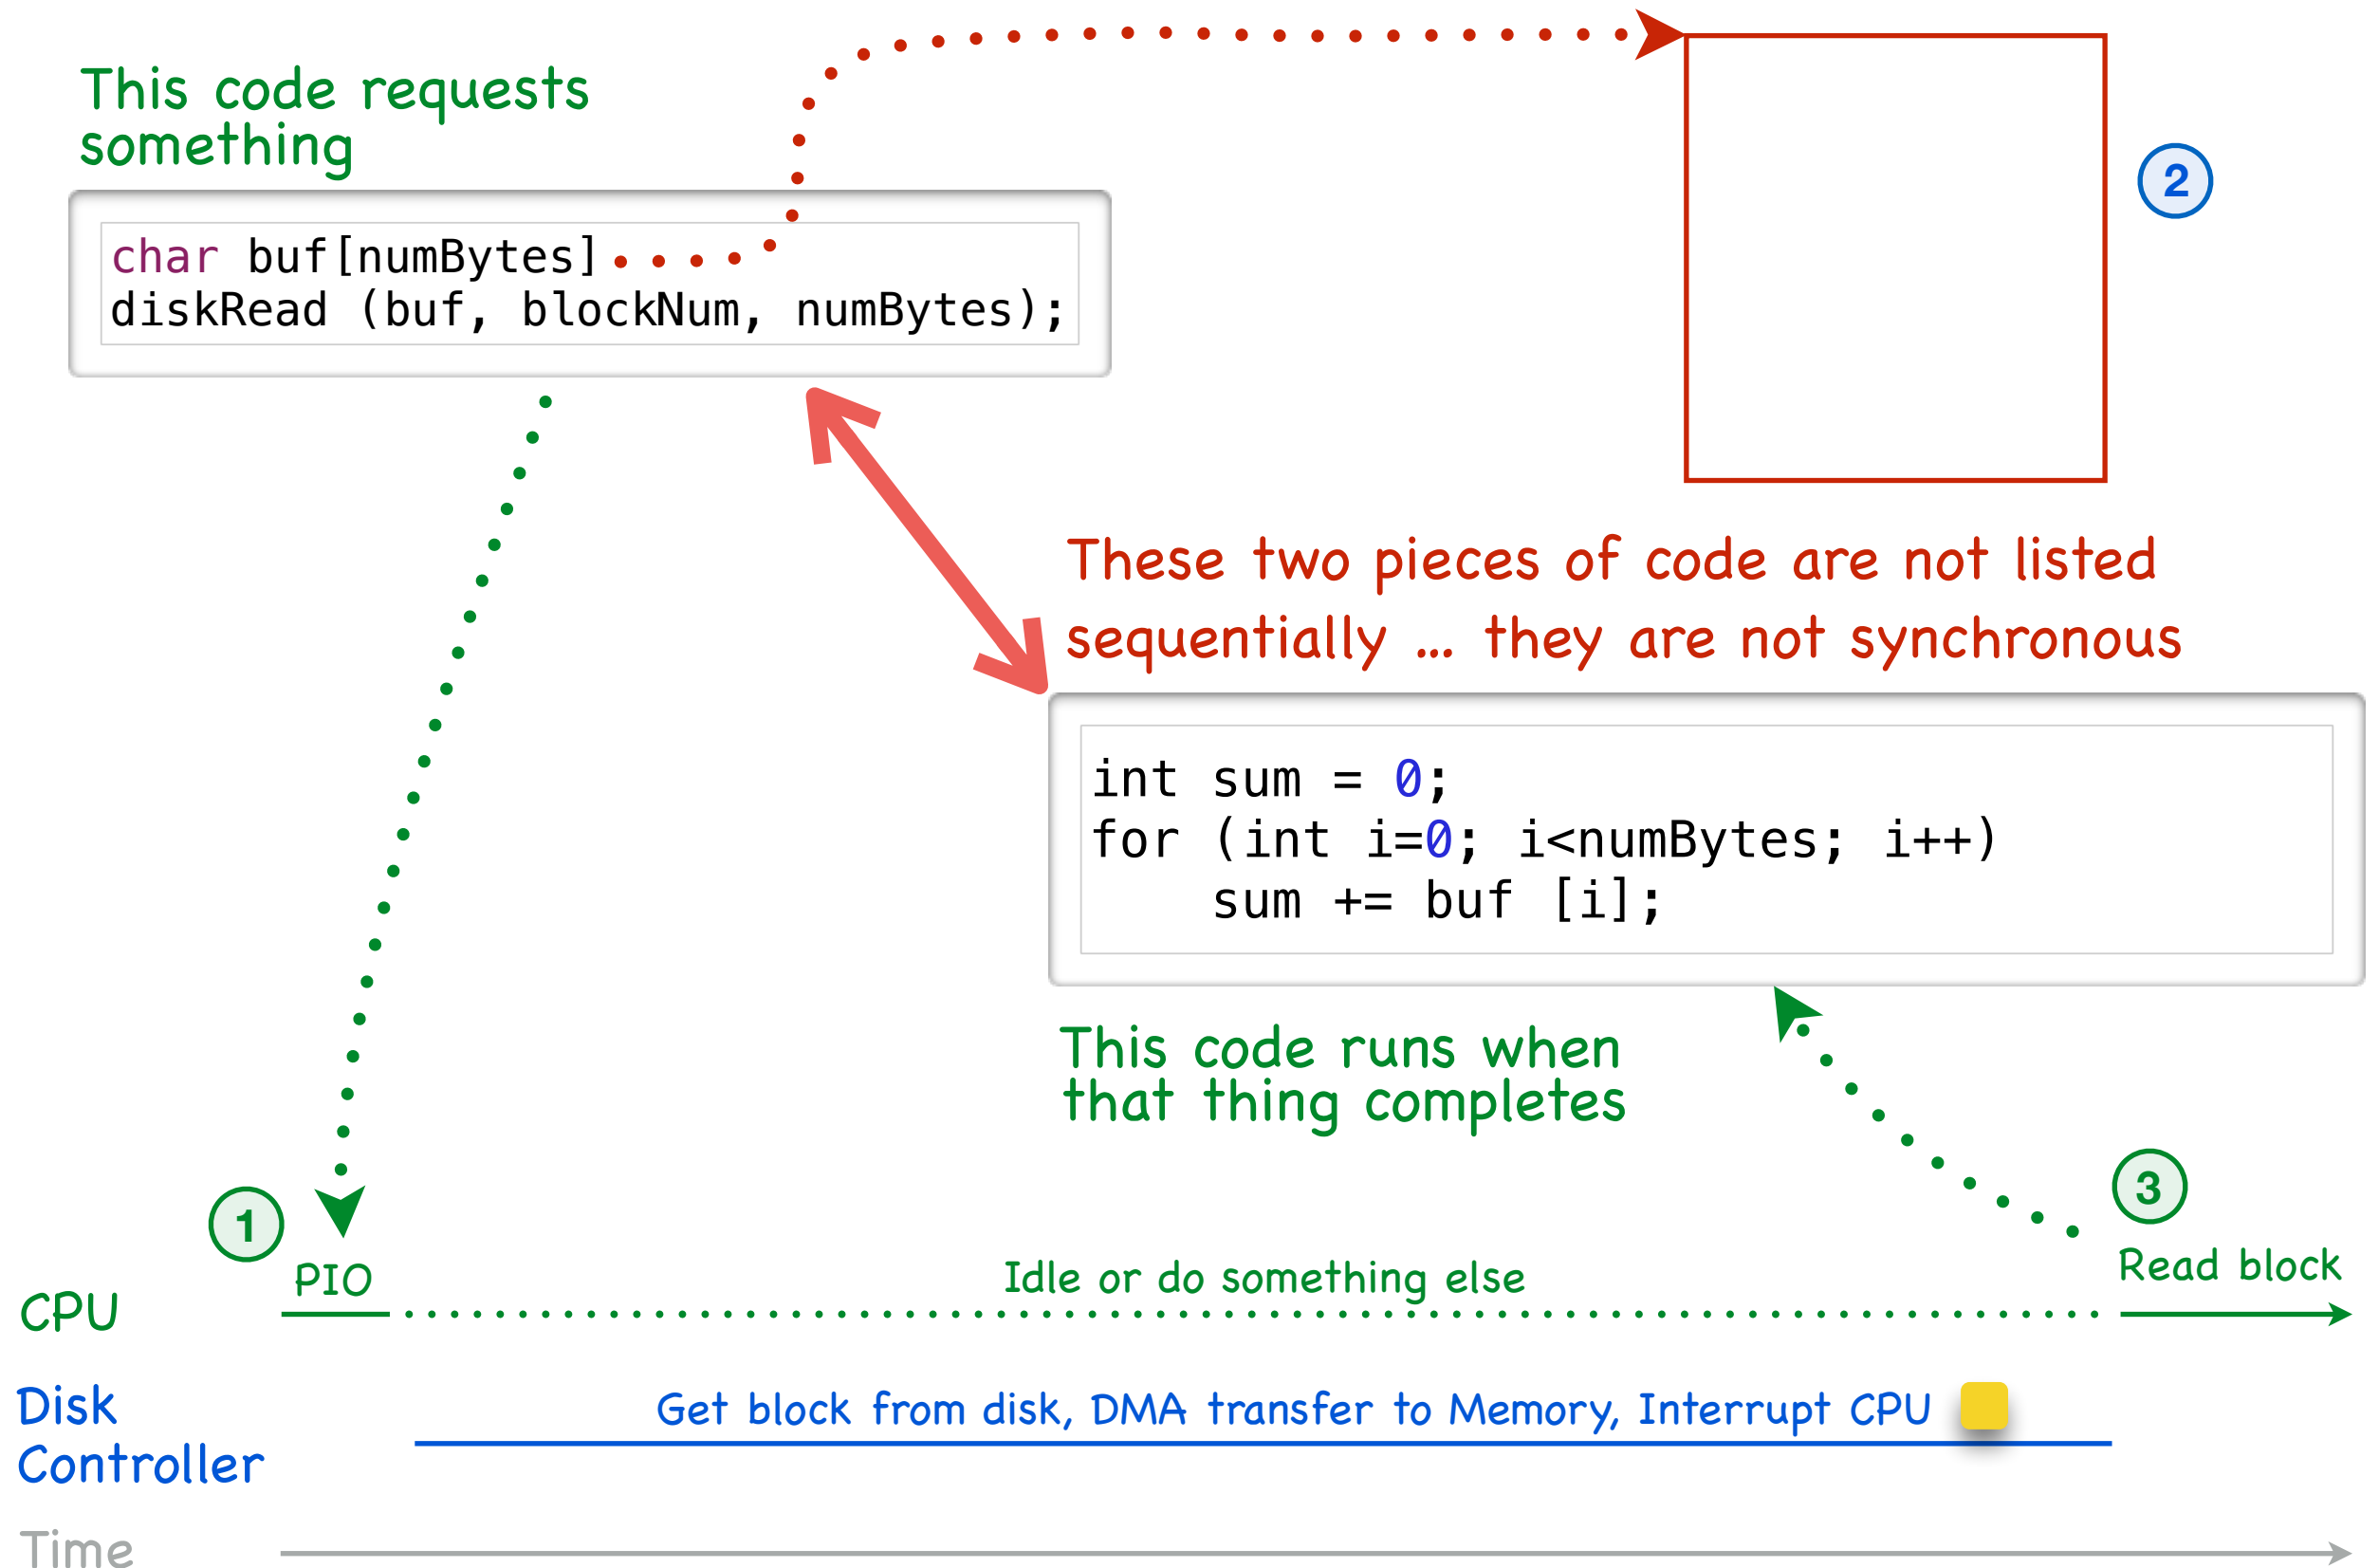

... Meets Reality

```
int sumDiskData(blockNum, numBytes) {  
    char buf[numBytes];  
    int sum = 0;  
  
    diskRead(buf, blockNum, numBytes);  
    for(int i=0; i<numBytes; i++)  
        sum += buf[i];  
    return sum;  
}
```

must wait for
read to complete



Making Code Asynchronous



Writing Asynchronous Code in C

► Events and Event Handlers

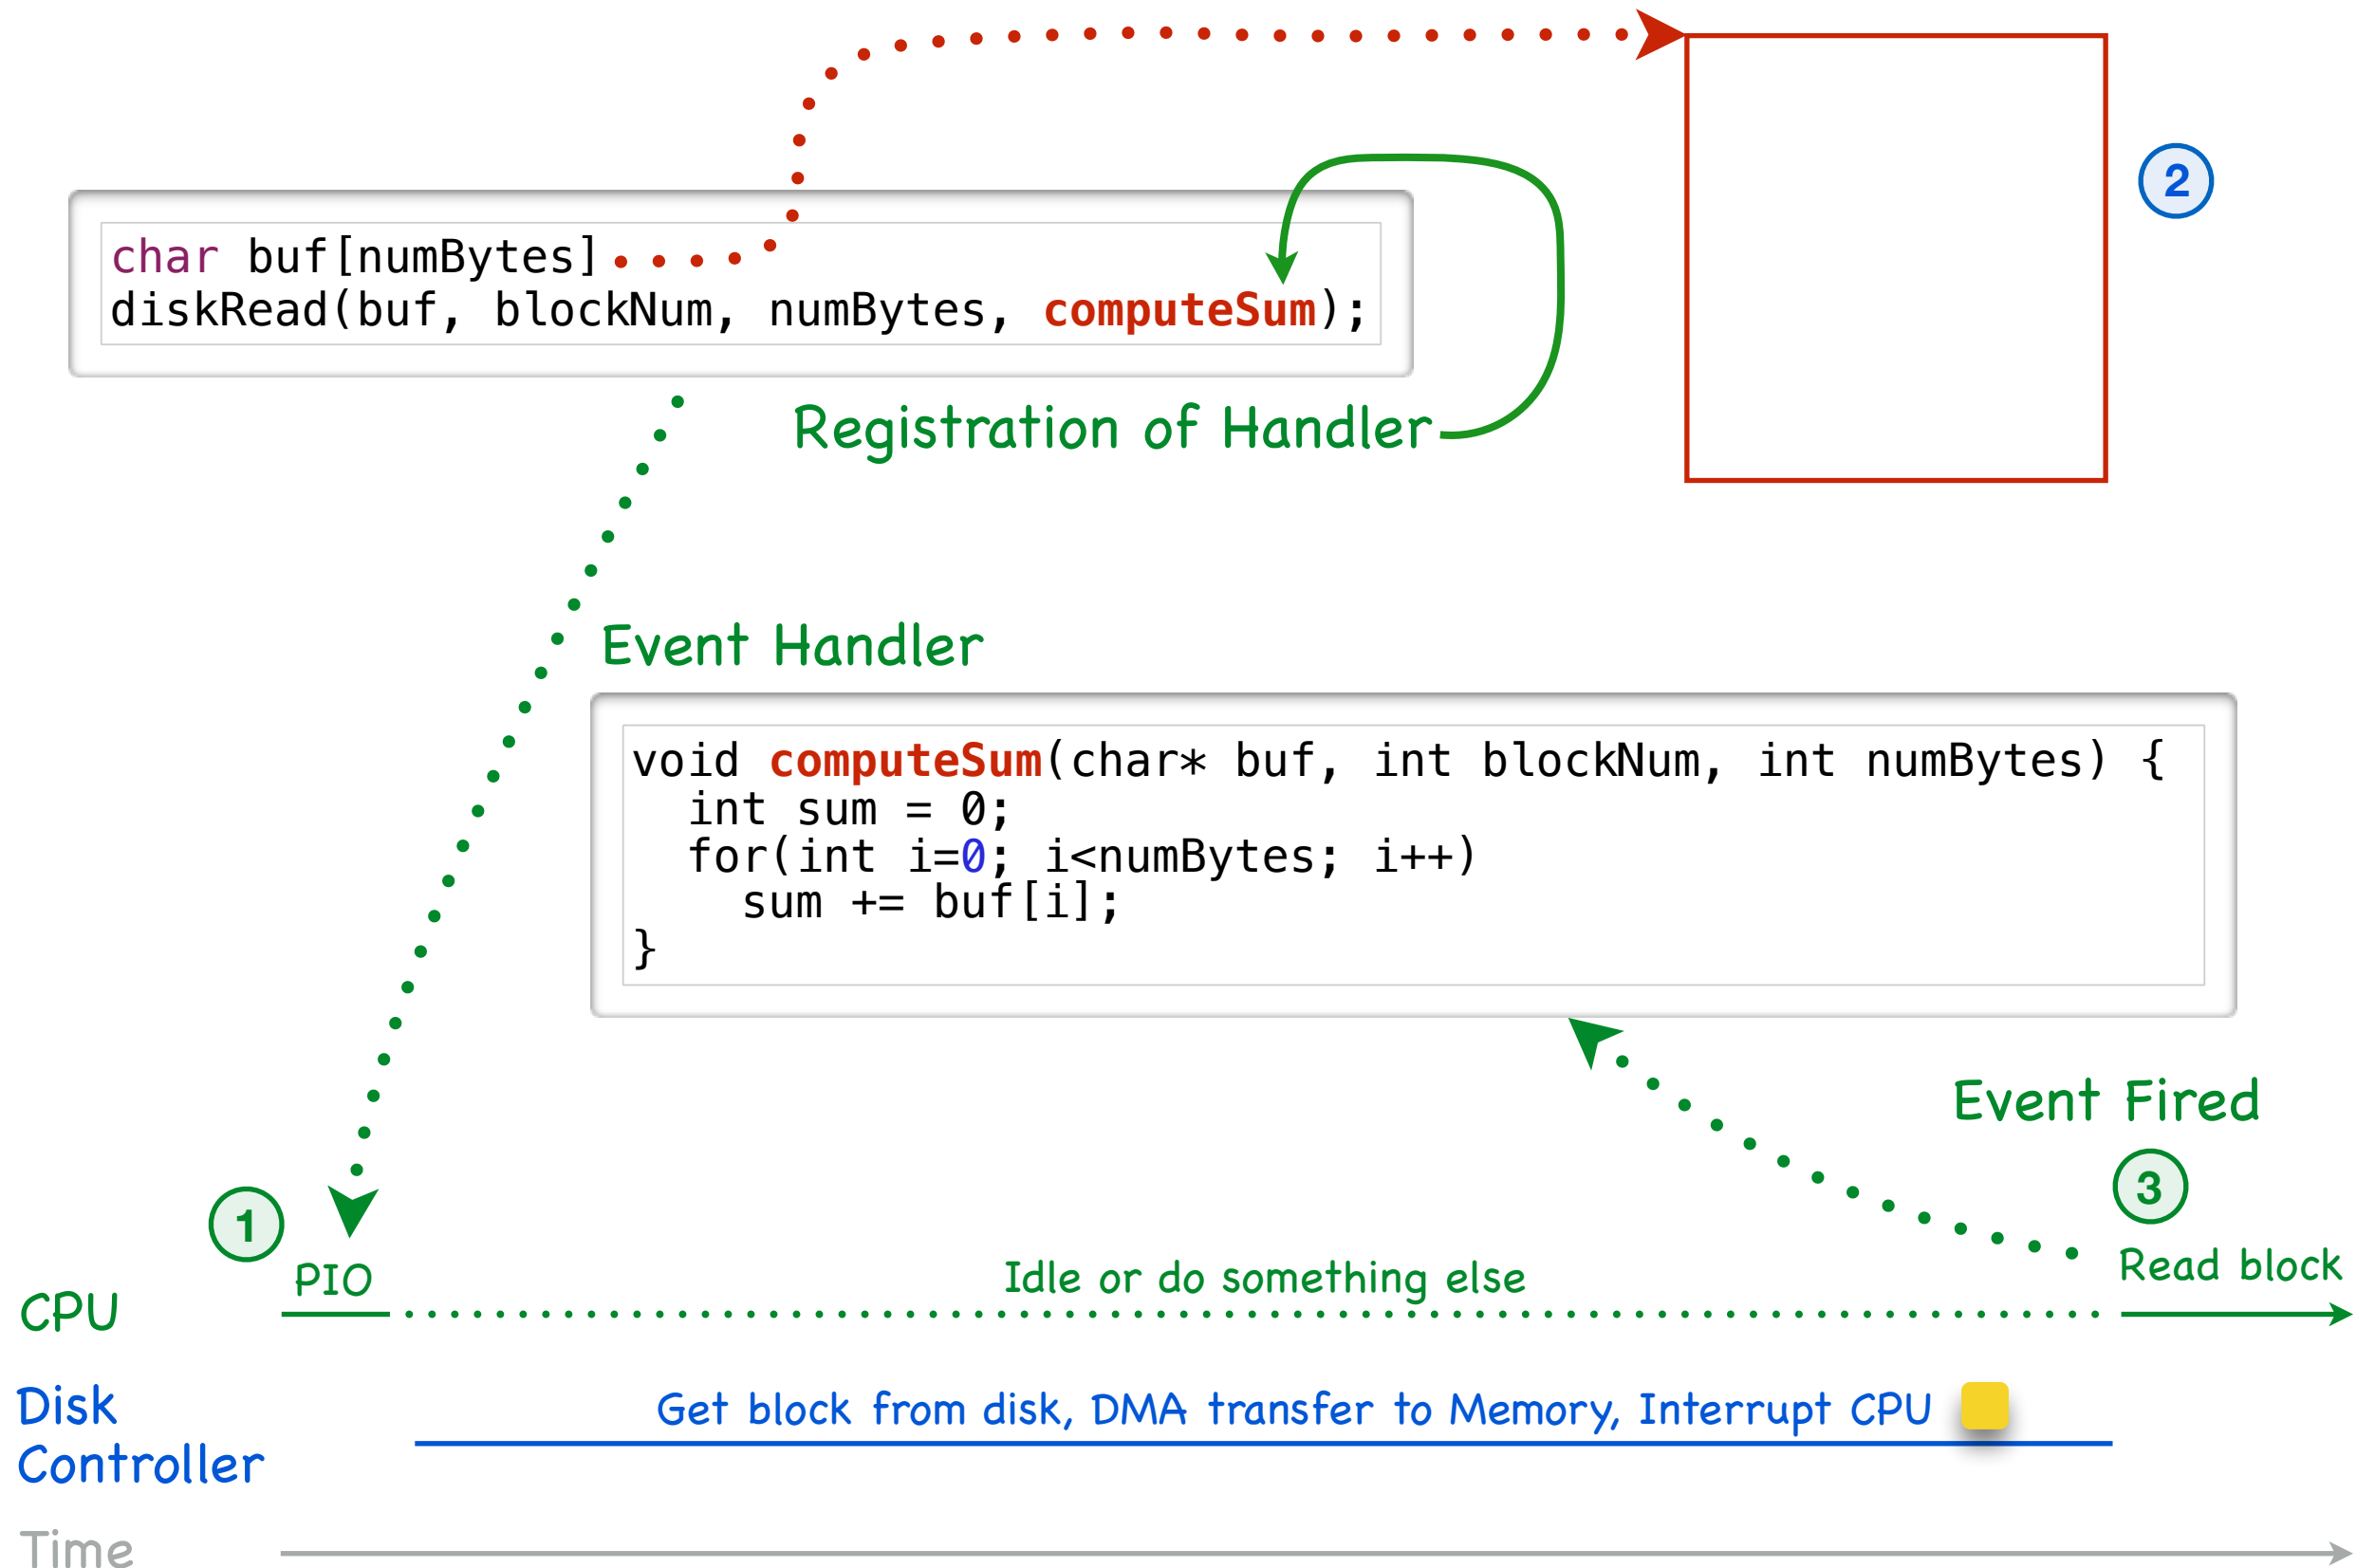
- the things that causes asynchronous code to run are called **events**
 - e.g., disk-read completion
- the code that runs when an event occurs is called an event **handler** (or *callback*)
 - e.g., the code that computes the checksum
- handlers are **registered** to a specific event
- events are **fired** to trigger the execution of the handler

► In Code

- request and registration of event handler are listed together in the code
- this code continues does not wait for event
- handler runs asynchronously when event occurs

```
void computeSum(char *buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for(int i=0; i<numBytes; i++)  
        sum += buf [i];  
}  
  
int sum(blockNum, numBytes) {  
    char buf[numBytes];  
  
    diskRead(buf, blockNum, numBytes, computeSum);  
}
```

Asynchronous Execution



Implementing Disk Read (Simplified)

▶ Interrupt Vector, Device ID, and PIO Address

- initialized before all this starts
- by operating system when device connects

```
#define MAX_DEVICES
void (*interruptVector [MAX_DEVICES])();

int  diskID      = 4;
int* diskAddress = (int*) 0x80001000;

interruptVector [diskID] = diskISR;
```

▶ Disk Read

- register event handler
- request block (PIO)

```
void diskRead(char *buf, int blockNum, int numBytes,
              void (*whenComplete)(char*, int, int))
{
    enqueue_handler(whenComplete, buf, blockNum, numBytes);
    // Perform PIO ... more in a moment
}
```

▶ Interrupt Handler

- find event handler and fire event
- firing means calling the handler procedure

```
void diskISR() {
    struct handler_dsc {
        void (*handler)(char*, int, int);
        char* buf;
        int  blockNum;
        int  numBytes;
    };
    struct handler_dsc hd;
    dequeue_handler (&hd);
    hd.handler (hd.buf, hd.blockNum, hd.numBytes);
}
```

Disk Read PIO

► Expanded Disk Read with PIO

- the message sent to disk has several fields
- each field had different device-memory address
- access it like a struct
 - but, keep in mind that writes are to device-memory;
 - they are messages across bus to device controller they are not writes to main memory

```
void *diskAddress = (void*) 0x80001000;
```

```
#DEFINE DISK_OP_READ  1  
#DEFINE DISK_OP_WRITE 2
```

```
struct disk_ctl {  
    int    op;  
    char *buf;  
    int    blockNum;  
    int    numBytes;  
}
```

```
void diskRead(char *buf, int blockNum, int numBytes, ...) {  
    struct disk_ctl *dc = diskAddress;  
    enqueue_handler (whenComplete, buf, blockNum, numBytes);  
    dc->op          = DISK_OP_READ;  
    dc->buf         = buf;  
    dc->blockNum    = blockNum;  
    dc->numBytes    = numBytes;  
}
```

Did We Really Solve The Problem?

- ▶ We wanted to do this

```
int sumDiskData(blockNum, numBytes) {  
    char buf[numBytes];  
    int sum = 0;  
  
    diskRead(buf, blockNum, numBytes);  
    for (int i=0; i<numBytes; i++)  
        sum += buf[i];  
  
    return sum;  
}
```

- ▶ But, reality forced us to do this

```
void computeSum(char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i=0; i<numBytes; i++)  
        sum += buf[i];  
}  
  
void sumDiskData(blockNum, numBytes) {  
    char buf[numBytes];  
  
    diskRead(buf, blockNum, numBytes, computeSum);  
}
```

- ▶ What's wrong?

Connecting Asynchrony to Program

- ▶ How do we use the value computed from the disk block
 - lets say we want to print it

```
void something() {  
    ...  
    int s = sumDiskData(buf, blk, n);  
    printf("%d\n", s);  
}
```

- but asynchronously?

```
void computeSum(char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i=0; i<numBytes; i++)  
        sum += buf[i];  
    free(buf);  
}  
  
void sumDiskData (blockNum, numBytes) {  
    char *buf = malloc(numBytes);  
  
    diskRead (buf, blockNum, numBytes, computeSum);  
}
```

Ordering in Asynchronous Programs

► If something has to happen after event

- it must be triggered by the event
- often this means it must be part of (or called by) event's handler

► To print the checksum

```
void computeSumAndPrint(char *buf, int blockNum, int numBytes) {
    int sum = 0;
    for (int i=0; i<numBytes; i++)
        sum += buf[i];
    printf("%d\n", sum);
    free(buf);
}

void sum(blockNum, numBytes, whenComplete) {
    char *buf = malloc(numBytes);

    diskRead(0, blockNum, numBytes, whenComplete);
}

sum (1234, 4096, computeSumAndPrint);
```

► Huge problem

- often there's a ton of stuff that depend on returned data
- that is, that must come after a particular event

Improving the Code ... But making it worse

```
void computeSumAndPrint(char *buf, int blockNum, int numBytes) {
    int sum = 0;
    for (int i=0; i<numBytes; i++)
        sum += buf[i];
    printf("%d\n", sum);
    free(buf);
}
```

```
void printInt (int i) {
    printf ("%d\n", i);
}

void computeSumAndCallback (... , void (*sumCallback) (int)) {
    int sum = 0;
    for (int i=0; i<numBytes; i++)
        sum += buf [i];
    sumCallback (sum);
    free (buf);
}

void sum (blockNum, numBytes, whenComplete, sumCallback) {
    char* buf = malloc (numBytes);

    diskRead (buf, blockNum, numBytes, whenComplete, sumCallback);
}

sum (1234, 4096, computeSumAndCallback, printInt);
```

What if you want
to do something
after printInt?

Welcome to
"callback hell" ...

Happy System, Sad Programmer

▶ Humans like synchrony

- we expect each step of a program to complete before the next one starts
- we use the result of previous steps as input to subsequent steps
- with disks, for example,
 - we read from a file in one step and then usually use the data we've read in the next step

▶ Computer systems are asynchronous

- the disk controller takes 10-20 milliseconds (10^{-3} s) to read a block
 - CPU can execute 60 million instructions while waiting for the disk to complete one read
 - we must allow the CPU to do other work while waiting for I/O completion
- many devices send unsolicited data at unpredictable times
 - e.g., incoming network packets, mouse clicks, keyboard-key presses
 - we must allow programs to be interrupted many, many times a second to handle these things

▶ Asynchrony makes programmers sad

- it makes programs more difficult to write and much more difficult to debug

Possible Solutions

▶ Accept the inevitable

- use an event-driven programming model
 - event triggering and handling are de-coupled
- a common idiom in many Java programs
 - GUI programming follows this model
- *CSP* (communicating sequential processes) boosts this idea to first-class status
 - no procedures or procedure calls
 - program code is decomposed into a set of sequential/synchronous processes
 - processes can fire events, which can cause other processes to run in parallel
 - each process has a guard predicate that lists events that will cause it to run
- Javascript in web browsers and Node.js embrace asynchrony, albeit awkwardly

▶ Invent a new abstraction

- an abstraction that provides programs the illusion of synchrony
- but, what happens when
 - a program does something asynchronous, like disk read?
 - an unanticipated device event occurs?

▶ What's the right solution?

- we still don't know — this is one of the most pressing questions we currently face