

CPSC 213

Introduction to Computer Systems

Summer Session 2019, Term 2

Unit 2d - Aug 8 - NOT ON THE FINAL

Virtual Memory

Overview

▶ Reading

- text: 9.1-9.2, 9.3.2-9.3.4

▶ Learning Goals

- explain how multiple programs can reside in memory concurrently, but preserve the single-system illusion
- explain why one program can not access the memory of another program and why this is important
- give key benefit and key drawback of base-and-bounds translation
- give key benefit and key drawback of segmentation
- give key benefit of paging
- translate a virtual address to a physical address using any of these three schemes

Multiple Concurrent Program Executions

▶ The last lie

- we've been assuming that the computer executes one program at a time
- but, really multiple programs can execute simultaneously

▶ Each Program Execution is

- called a ***process***
- multiple threads executing a program with state stored in memory
- compiler-assigned addresses for static things (e.g., code and globals)
- memory is viewable and changeable only by the program execution

▶ Problem

- there is only one memory that they must all share
- compiler-assigned addresses overlap
 - its okay for multiple processes to share code (or anything that is read-only)
 - global variables must be private to each process, but they can be at the same address
- programs can access any address
 - must ensure that each process's memory is completely isolated

Virtual Memory

► Basic Idea


- each process is allocated a portion of physical memory
- allowed to use its own addresses for it
- disallowed from accessing any other part of memory

► Changes to Program Execution

- hardware translates address of every memory access
 - by *memory management unit (MMU)*: part of CPU
- from process's *virtual* address to machine's real *physical* address

► Virtual Address Space

- one per process
- maps process's virtual addresses to machine's physical addresses
- mapping data structure
 - created and maintained by operating system
 - accessed directly by hardware to translate addresses on every memory access



Implementing the MMU

- ▶ Lets think of this in the simulator ...
 - introduce a class to simulate the MMU hardware

```
class MMU extends MainMemory {  
    byte []      physicalMemory;  
    AddressSpace currentAddressSpace;  
  
    void setAddressSpace (AddressSpace* as);  
  
    byte readByte (int va) {  
        int pa = currentAddressSpace.translate (va);  
        return physicalMemory.read (pa);  
    }  
}
```

- currentAddressSpace is a hardware register
- the address space performs virtual-to-physical address translation

Implementing Address Translation

```
class MMU extends MainMemory {  
    byte []      physicalMemory;  
    AddressSpace currentAddressSpace;  
  
    void setAddressSpace (AddressSpace* as);  
  
    byte readByte (int va) {  
        int pa = currentAddressSpace.translate (va);  
        return physicalMemory.read (pa);  
    }  
}
```

► Goal

- translate any virtual address to a unique physical address (or none)
- fast and efficient hardware implementation

► Lets look at a couple of alternatives ...

Base and Bounds


► An address space is

- a single, variable-size, non-expandable chunk of physical memory
- named by its base physical address and its length

► As a class in the simulator

```
class AddressSpace {  
    int baseVA, basePA, bounds;  
  
    int translate (int va) {  
        int offset = va - baseVA;  
        if (offset < 0 || offset > bounds)  
            throw new IllegalArgumentException ();  
        return basePA + offset;  
    }  
}
```

► Problems



But, Address Space Use May Be Sparse

► Issue


- the address space of a program execution is divided into regions
- for example: code, globals, heap, shared-libraries and stack
- there are large gaps of unused address space between these regions

► Problem

- a single base-and-bounds mapping from virtual to physical addresses
- means that gaps in virtual address space will waste physical memory
- this is the **Internal Fragmentation** problem



► Solution



Segmentation

► Address Space

- a set of segments

► Segment


- a single, variable-size, non-expandable chunk of physical memory
- named by its base virtual address, physical address and length

► Implementation in Simulator

```
class AddressSpace {
    Segment segment[];

    int translate (int va) {
        for (int i=0; i<segments.length; i++) {
            int offset = va - segment[i].baseVA;
            if (offset >= 0 && offset < segment[i].bounds) {
                pa = segment[i].basePA + offset;
                return pa;
            }
        }
        throw new IllegalArgumentException (va);
    }
}
```

► Problem




Memory Use is Not Known Staticly

► Issue

- segments are **not expandable**; their size is fixed when segment is created
- some segments such as stack and heap change size dynamically

► Problem

- segment size is chosen when segment is created
- too large and internal fragmentation wastes memory
- too small and stack or heap restricted



► Solution

- allow segments to expand?


There May Be No Room to Expand

► Issue

- segments are contiguous chunks of physical memory
- a segment can only expand to fill space between it and the next segment

► Problem

- there is no guarantee there will be room to expand a segment
- the available memory space is not where we want it (i.e., adjacent to segment)
- this is the **External Fragmentation** problem



► Solution


Moving Segments is Expensive

► Issue

- if there is space in memory to store expanding segment, but not where it is
- could move expanding segment or other segments to make room
- external fragmentation is resolved by moving things to consolidate free space

► Problem

- moving is possible, but expensive
- to move a segment, all of its data must be copied
- segments are large and memory copying is expensive




Expand Segments by Adding Segments

▶ What we know

- segments should be non-expandable
- size can not be effectively determined statically

▶ Idea

- instead of expanding a segment
- make a new one that is adjacent virtually, but not physically




▶ Problem

- oh no! another problem! what is it? why does it occur?

Eliminating External Fragmentation

- ▶ The problem with what we are doing is
 - allocating variable size segments leads to external fragmentation of memory
 - this is an inherent problem with variable-size allocation
- ▶ What about **fixed sized allocation**
 - could we make every segment the same size?
 - this eliminates external fragmentation
 - but, if we make segments too big, we'll get internal fragmentation
 - so, they need to be fairly small and so we'll have lots of them



▶ Problem

Translation with Many Segments

- ▶ What is wrong with this approach if there are many segments?

```
class AddressSpace {
    Segment segment[];

    int translate (int va) {
        for (int i=0; i<segment.length; i++) {
            int offset = va - segment[i].baseVA;
            if (offset > 0 && offset < segment[i].bounds) {
                pa = segment[i].basePA + offset;
                return pa;
            }
        }
        throw new IllegalArgumentException (va);
    }
}
```


- ▶ Now what?

- is there another way to locate the segment, when segments are fixed size?

Paging

▶ Key Idea

- Address Space is divided into set of fixed-size segments called pages
- every page of virtual address space has a number (**VPN**)
- virtual address (**VA**) is comprised of
 - $\text{VPN} = \text{VA} / \text{PAGE_SIZE}$
 - $\text{OFFSET} = \text{VA} \% \text{PAGE_SIZE}$
- every page *frame* of physical memory has a number too (**PFN**) that yields physical address (**PA**)
 - $\text{PA} = \text{PFN} * \text{PAGE_SIZE} + \text{OFFSET}$



▶ Page Table


- stored in memory
 - (address with physical addresses or operating-system's virtual addresses)
- one for every address space
- array of page-table entries, one for virtual page in address space
- indexed by VPN

▶ Page Table Entry

- indicates whether virtual page is valid
- if valid stores the page's PFN
- also stores access permissions
 - is reading okay? is writing okay? is executing okay?
- stores a bit of other stuff too (e.g., access and dirty flags)

► New terminology

- **page** a small, fixed-sized (4-KB) segment
- **page table** virtual-to-physical translation table
- **pte** page table entry
- **vpn** virtual page number
- **pfn** physical page frame number
- **offset** byte offset of address from beginning of page



► Translation using a Page Table

```
class PageTableEntry {  
    boolean isValid;  
    boolean isReadAllowed;  
    boolean isWriteAllowed;  
    boolean isExecuteAllowed;  
    int pfn;  
}
```

```
class AddressSpace {  
    PageTableEntry pte[];  
  
    int translate (int va) {  
        int vpn      = va / PAGE_SIZE;  
        int offset = va % PAGE_SIZE;  
        if (pte[vpn].isValid)  
            return pte[vpn].pfn * PAGE_SIZE + offset;  
        else  
            throw new IllegalArgumentException (va);  
    }  
}
```

▶ The bit-shifty version

- assume that page size is 4-KB = 4096 = 2^{12}
- assume addresses are 32 bits
- then, vpn and pfn are 20 bits and offset is 12 bits
- pte is pfn plus valid bit, so 21 bits or so, say 4 bytes
- page table has 2^{20} pte's and so is 4-MB in size

▶ The simulator code

```
class PageTableEntry {  
    boolean isValid;  
    int      pfn;  
}
```

```
class AddressSpace {  
    PageTableEntry pte[];  
  
    int translate (int va) {  
        int vpn      = va >>> 12;  
        int offset = va & 0xfff;  
        if (pte[vpn].isValid)  
            return pte[vpn].pfn << 12 | offset;  
        else  
            throw new IllegalArgumentException (va);  
    }  
}
```

Question

- ▶ Consider this page table



```
0x00000000  
0x80000007  
0x80000321  
0x8000006b  
0x8000005a  
0x80000040  
0x00000000
```

- ▶ Is 0x43a0 a valid virtual address and if so what is the corresponding physical address?
 - (A) Not valid
 - (B) 0x43a0
 - (C) 0x5a3a0
 - (D) 0x73a0
 - (E) 0x3a0

Translation and Exceptions

▶ Virtual-to-Physical translation

- occurs on every memory reference
- handled by hardware (sometimes with some software)
- aided by a cache of recent translations
- but, in general requires reading page table entry from memory

▶ Page fault

- is an exception raised by the CPU
- when a virtual address is invalid
- an exception is just like an interrupt, but generated by CPU not IO device
- page fault handler runs each time a page fault occurs

▶ Handling a page fault

- extending the heap or stack, handler can just deliver a new zero-filled page
- what about the code, global variables, or existing parts of heap or stack?

Demand Paging

▶ Key Idea

- some application data is not in memory
- transfer from disk to memory, only when needed

▶ Page Table


- only stores entries for pages that are in memory
- pages that are only on disk are marked invalid
- access to non-resident page- causes a page-fault interrupt

▶ Memory Map

- a second data structure managed by the OS
- divides virtual address space into regions, each mapped to a file
- page-fault interrupt handler checks to see if faulted page is mapped
- if so, gets page from disk, update Page Table and restart faulted instruction

▶ Page Replacement

- pages can now be removed from memory, transparent to program
- a replacement algorithm choose which pages should be resident and swaps out others



Context Switch

▶ A context switch is

- switching between threads from different processes
- each process has a private address space and thus its own page table

▶ Implementing a context switch

- change PTBR to point to new process's page table
- switch threads (save regs, switch stacks, restore regs)

▶ Context Switch vs Thread Switch

- changing page tables can be much slower than just changing threads
- mainly because caching techniques used to make translation fast

Summary

▶ Process

- a program execution
- a private virtual address space and a set of threads
- private address space required for static address allocation and isolation

▶ Virtual Address Space

- a mapping from virtual addresses to physical memory addresses
- programs use virtual addresses
- the MMU translates them to physical address used by the memory hardware

▶ Paging

- a way to implement address space translation
- divide virtual address space into small, fixed sized virtual page frames
- page table stores base physical address of every virtual page frame
- page table is indexed by virtual page frame number
- some virtual page frames have no physical page mapping
- some of these get data on demand from disk

Address Space Translation Tradeoffs

- ▶ **Single, variable-size, non-expandable segment**
 - internal fragmentation of segment due to sparse address use
- ▶ **Multiple, variable-size, non-expandable segments**
 - internal fragmentation of segments when size isn't known statically
 - external fragmentation of memory because segments are variable size
 - moving segments would resolve fragmentation, but moving is costly
- ▶ **Expandable segments**
 - expansion must be physically contiguous, but there may not be room
 - external fragmentation of memory requires moving segments to make room
- ▶ **Multiple, fixed-size, non-expandable segments**
 - called pages
 - need to be small to avoid internal fragmentation, so there are many of them
 - since there are many, need indexed lookup instead of search